

# Speech Recognition and Deep Learning

Adam Coates

Vinay Rao



# Speech recognition

- Many exciting & valuable applications



Content captioning



Cars / Hands-free interfaces



Home devices

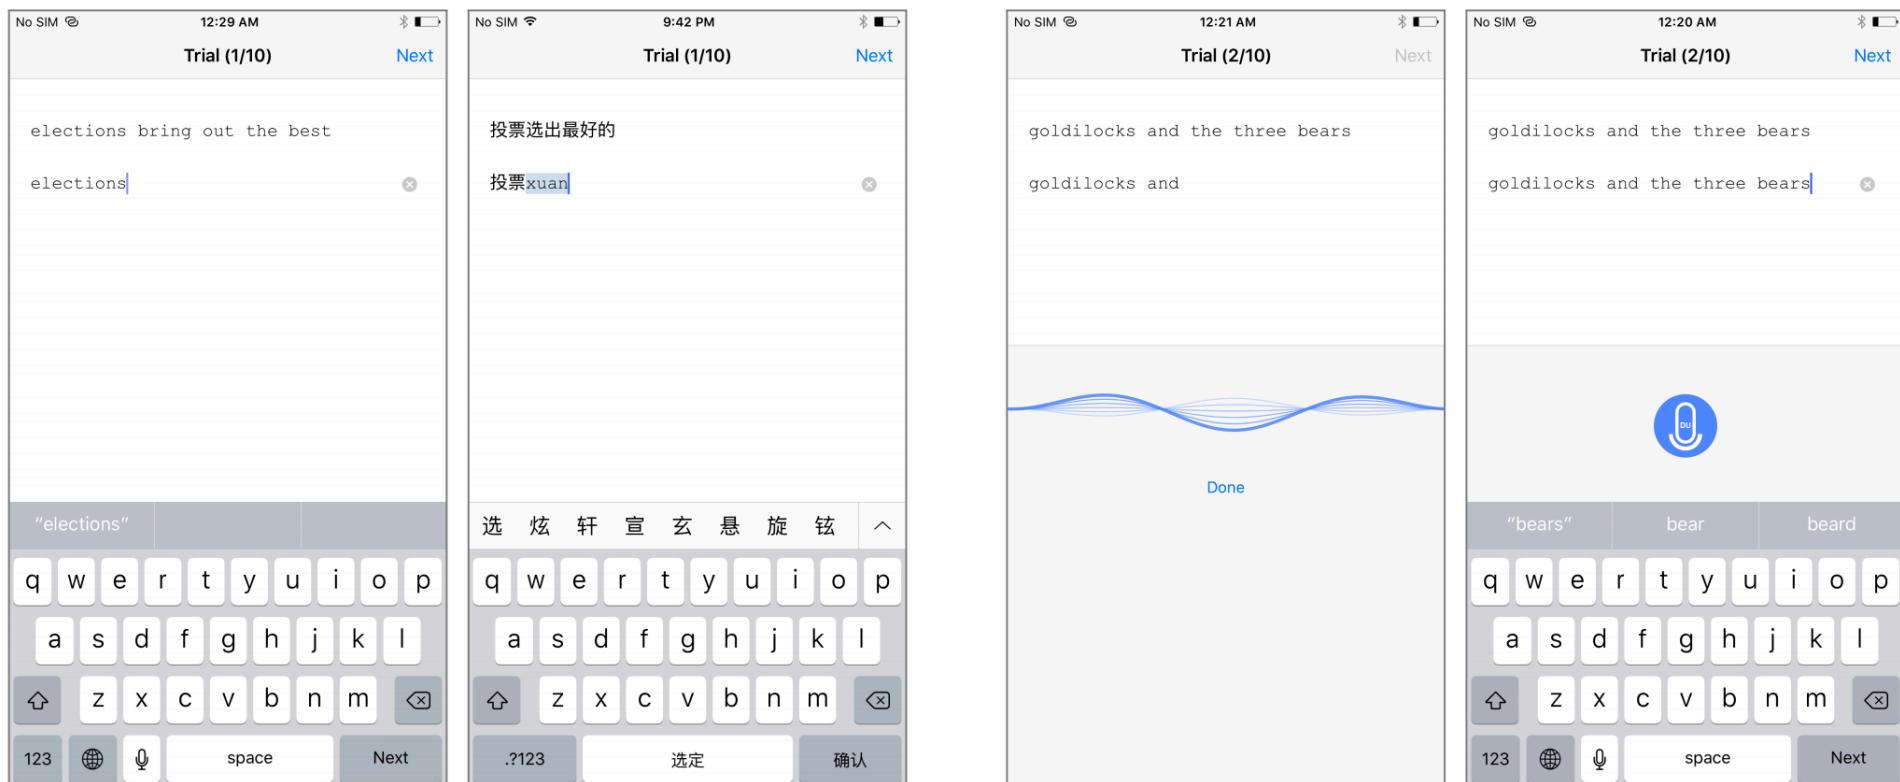


Mobile devices

# Speech recognition

- Many exciting & valuable applications

“Speech Is 3x Faster than Typing for English and Mandarin Text Entry on Mobile Devices”



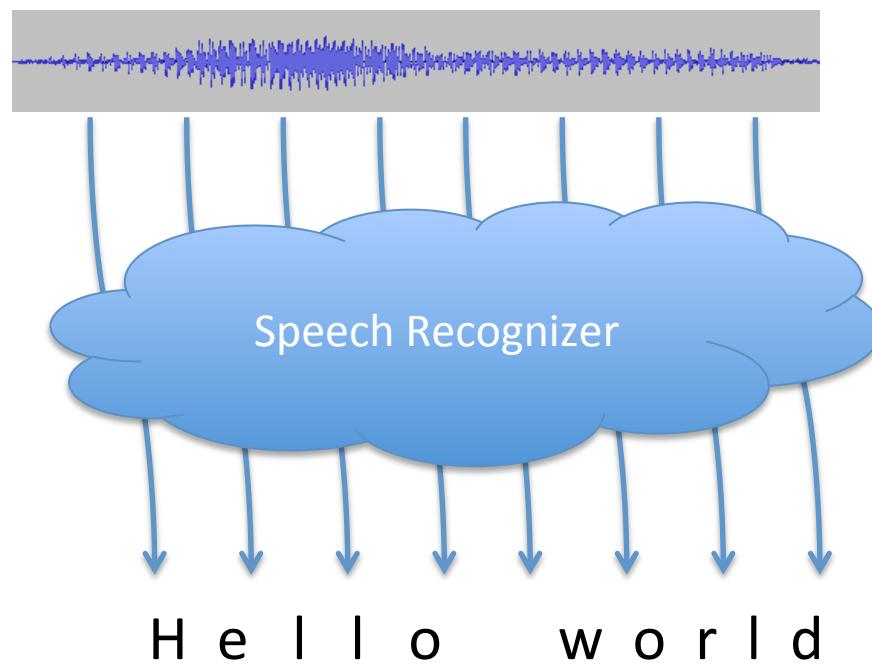
[Ruan et al., 2016]

# Speech recognition

- Many components make up a complete speech application:
  - Speech transcription This lecture
  - Word spotting / trigger word
  - Speaker identification / verification

# Speech recognition

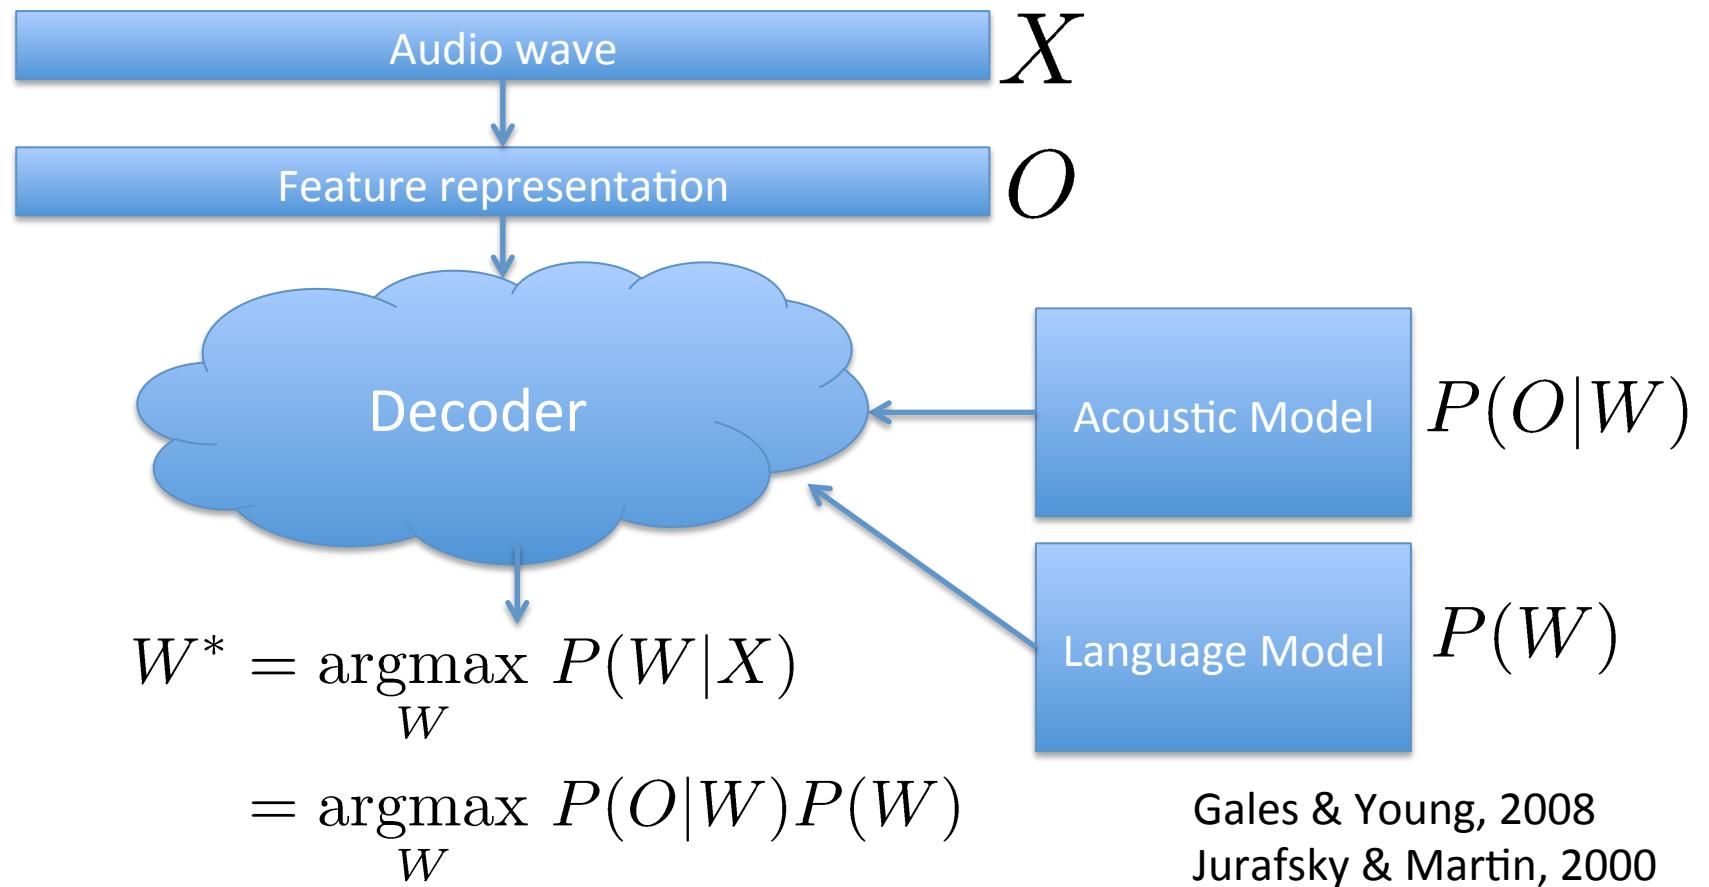
- Given speech audio, generate a transcript.



Important goal of AI: historically hard for machines, easy for people.

# Traditional ASR pipeline

- Traditional systems break problem into several key components:



# Traditional ASR pipeline

- Usually represent words as sequence of “phonemes”:

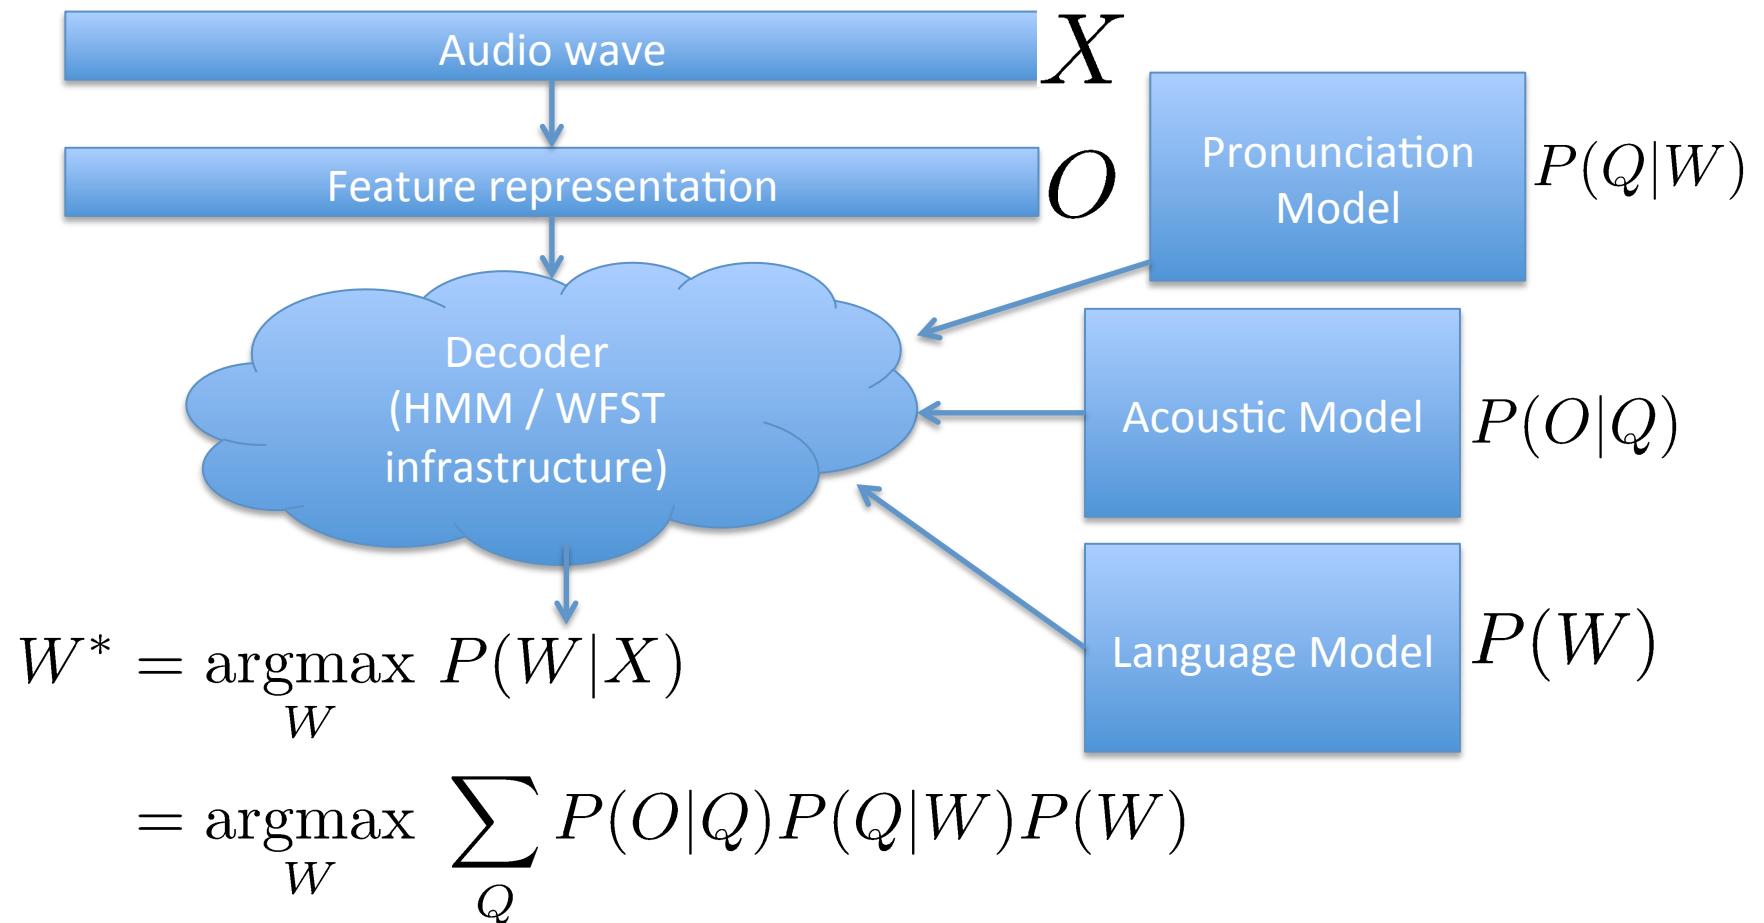
$$w_1 = \text{“hello”} = [\text{HH AH L OW}] = [q_1 q_2 q_3 q_4]$$

- Phonemes are the perceptually distinct units of sound that distinguish words.
  - Quite approximate... but sorta standardized-ish.
  - Some labeled corpora available (e.g., TIMIT)

	Phone Label	Example		Phone Label	Example		Phone Label	Example
1	iy	beet	22	ch	choke	43	en	button
2	ih	bit	23	b	bee	44	eng	Washington

# Traditional ASR pipeline

- Traditional systems usually model phoneme sequences instead of words. This necessitates a dictionary or other model to translate.



# Traditional ASR pipeline

- Traditional pipeline is highly tweak-able, but also hard to get working well.
- Historically, each part of system has own set of challenges.
  - E.g., choosing feature representation.

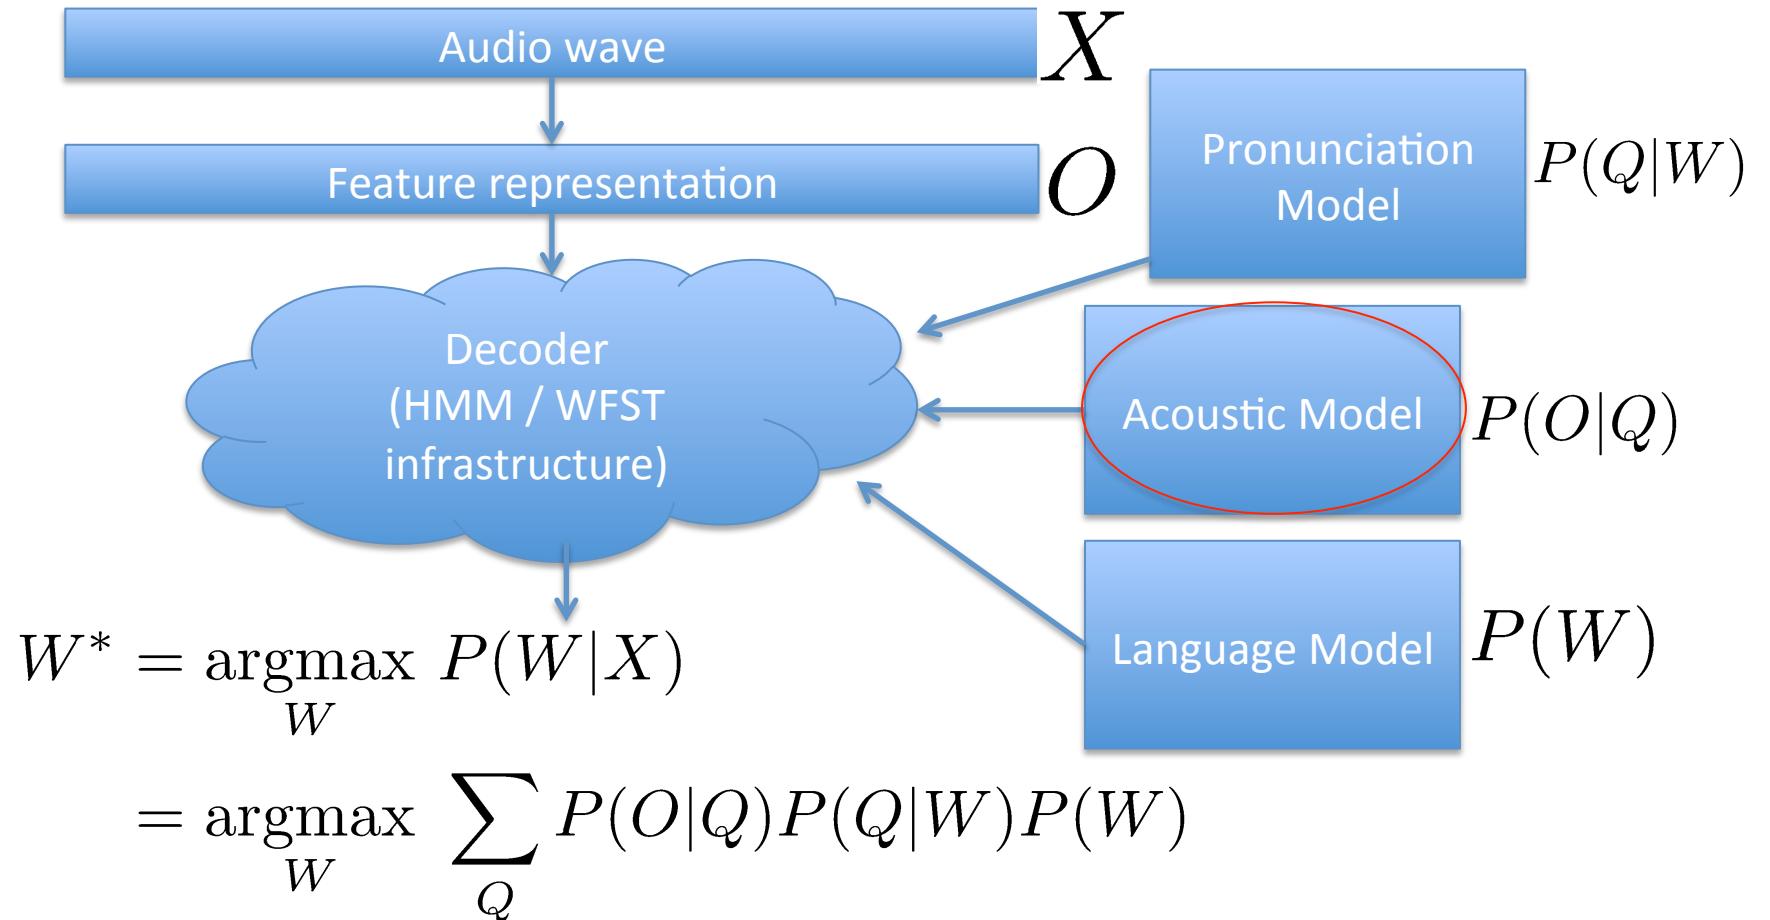
# Deep Learning in ASR

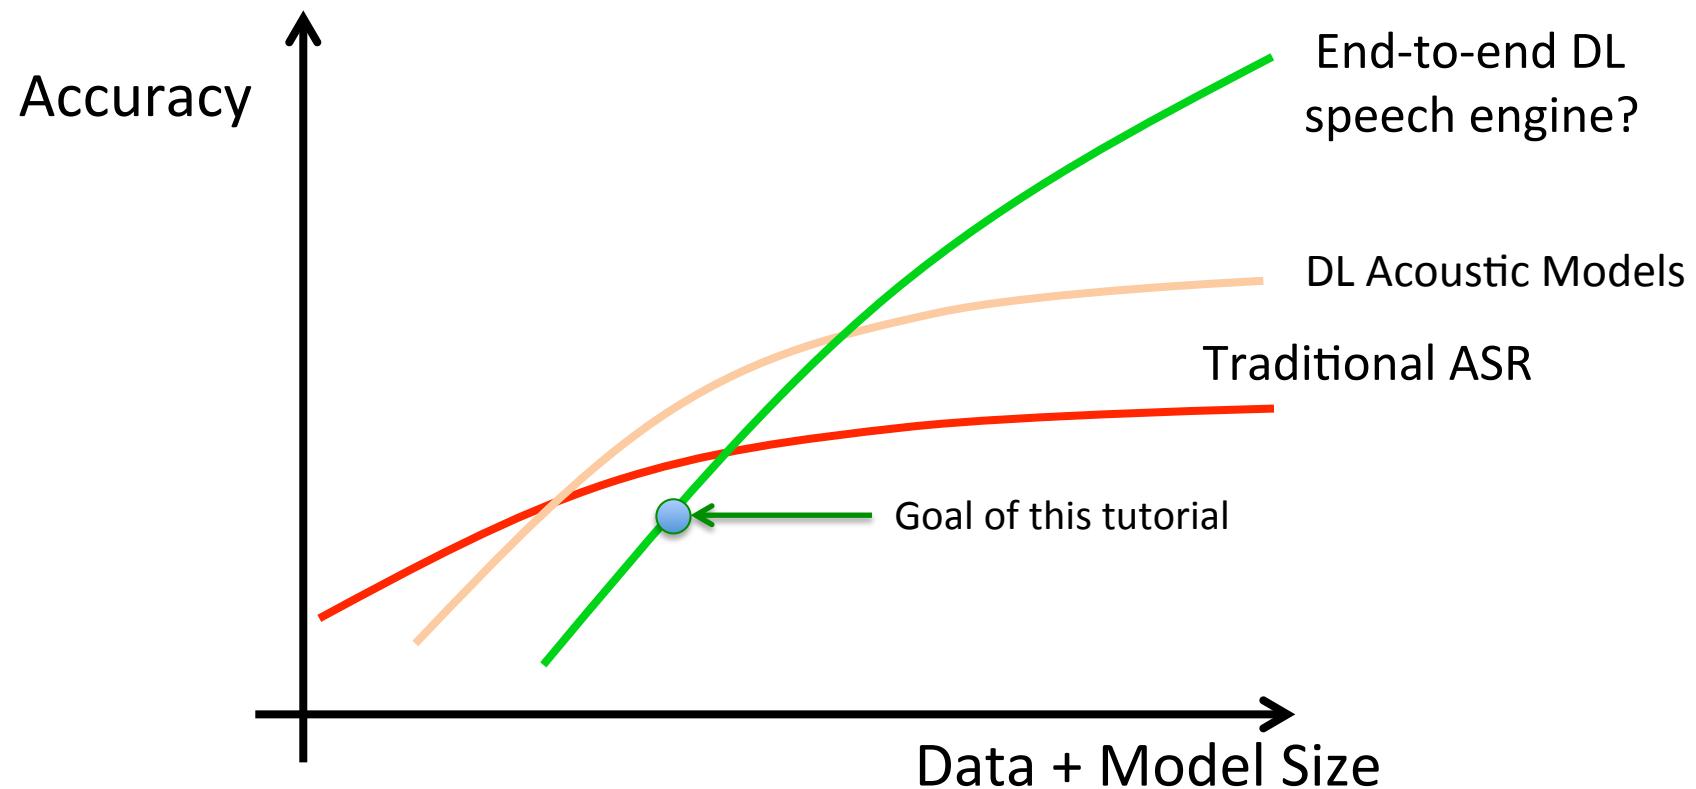
- Where to apply DL to make ASR better?
  - Good start: improve acoustic model  $P(O|Q)$
- Introduction of pre-training/DBN:

DBN-HMM	5	from DBN-HMM	Triphone Senones	yes	71.8%	69.6%
ML GMM-HMM baseline					62.9%	60.4%
MMI GMM-HMM baseline					65.1%	62.8%
MPE GMM-HMM baseline					65.5%	63.8%
ML GMM-HMM baseline 2100 hours of training data (transcription is 90% accurate)				-	62.9%	[13]

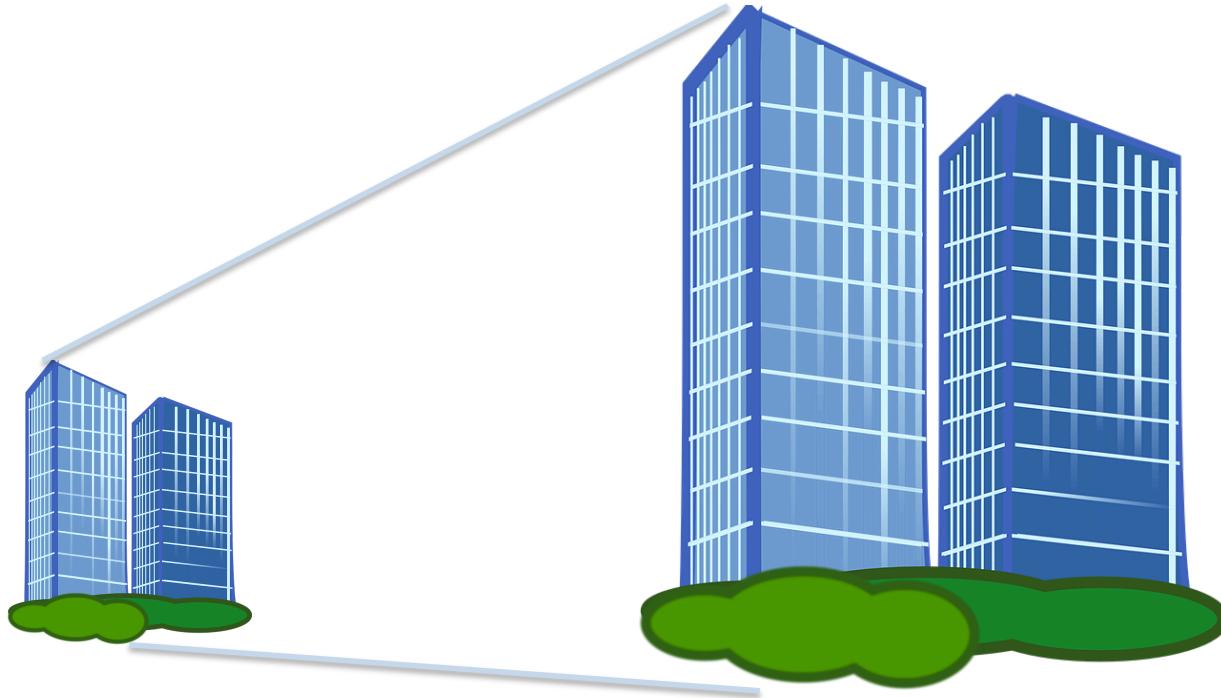
[Dahl et al., ICASSP 2011]

# Traditional ASR pipeline





# Scale model



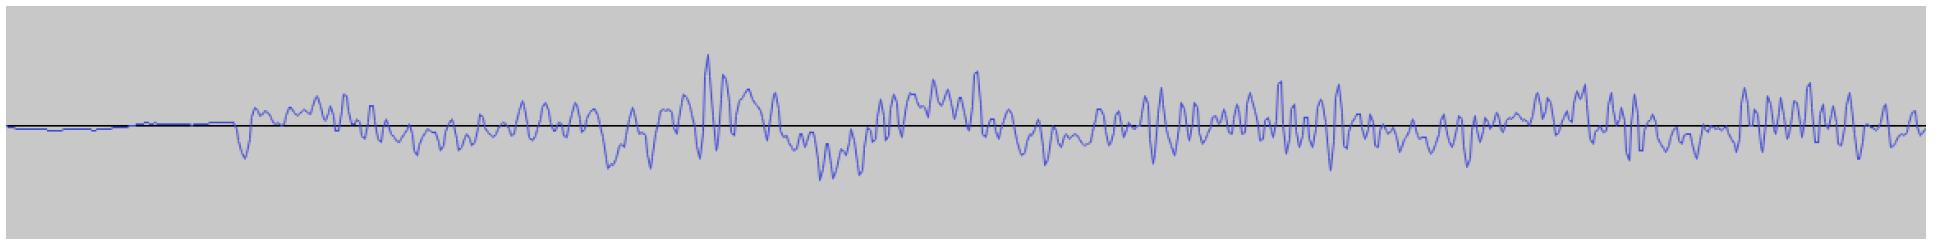
Starter code: [github.com/baidu-research/ba-dls-deepspeech](https://github.com/baidu-research/ba-dls-deepspeech)  
Get a simple max-decoded pipeline running.

# Outline

- DL speech pipeline walkthrough
  - Preprocessing
  - CTC
  - Training
  - Decoding & language models
- Scaling up!
  - Data
  - Systems
- Production / reality

# Raw audio

- Simple 1D signal:



Typical sample rates for speech: 8KHz, 16KHz.

Each sample typically 8-bit or 16-bit.

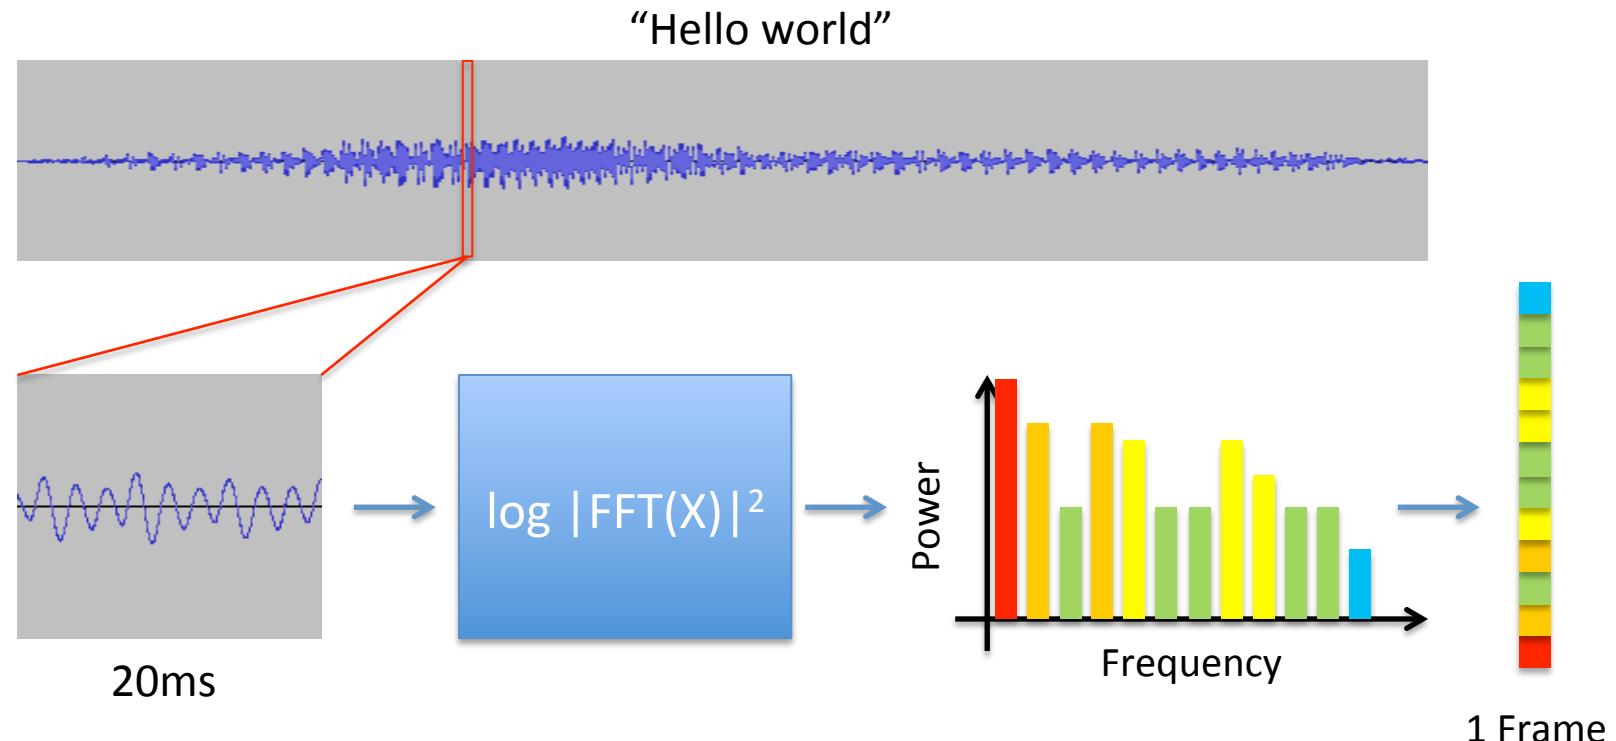
- 1D vector:  $X = [x_1 x_2 \dots]$

# Pre-processing

- Two ways to start:
    - Minimally pre-process (e.g., simple spectrogram).
      - We'll use this.
    - Train model from raw audio wave.
      - It works!
- See, e.g., Sainath et al., Interspeech 2015

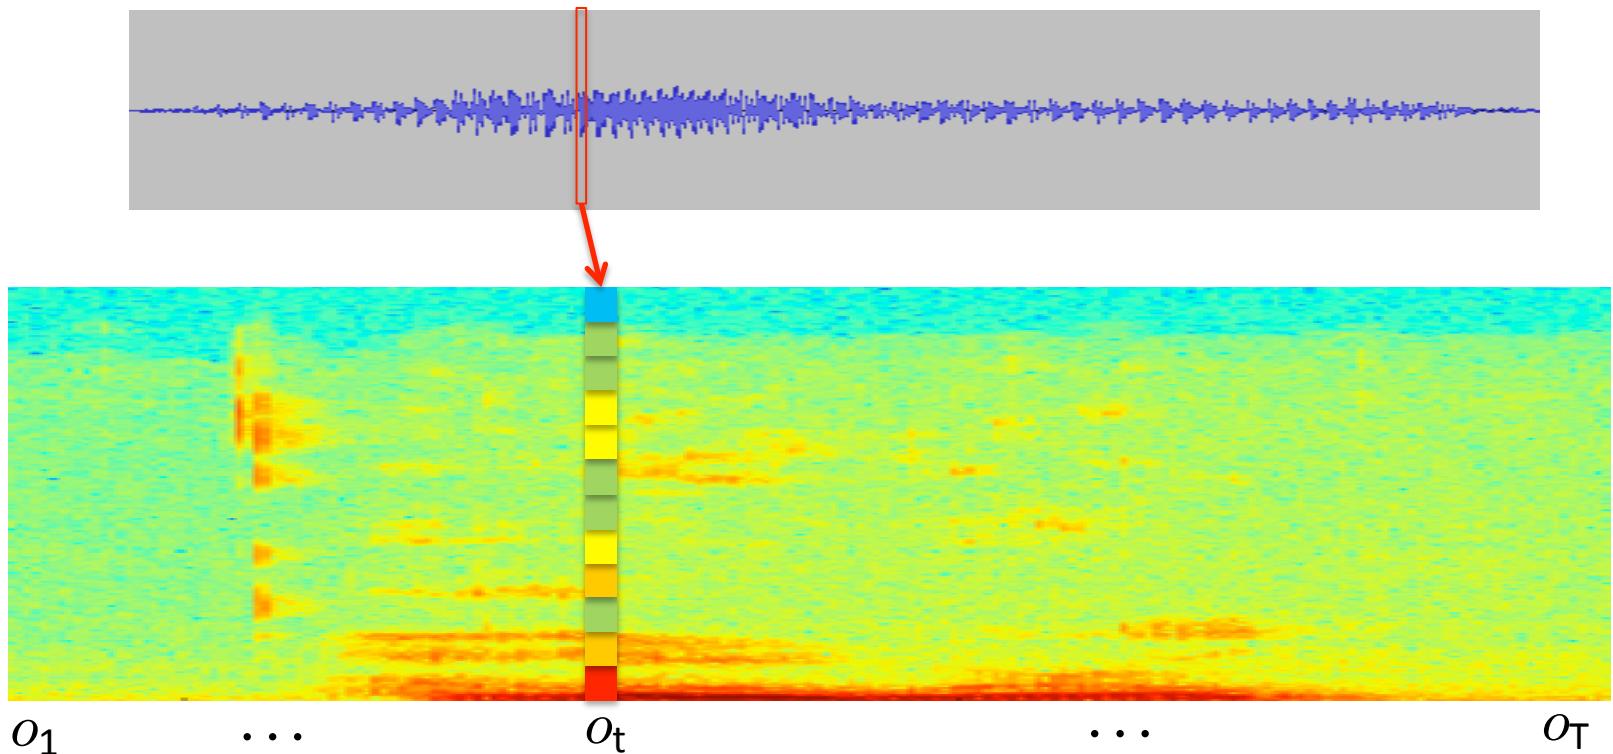
# Spectrogram

- Take a small window (e.g., 20ms) of waveform.
  - Compute FFT and take magnitude. (i.e., power)
  - Describes frequency content in local window.



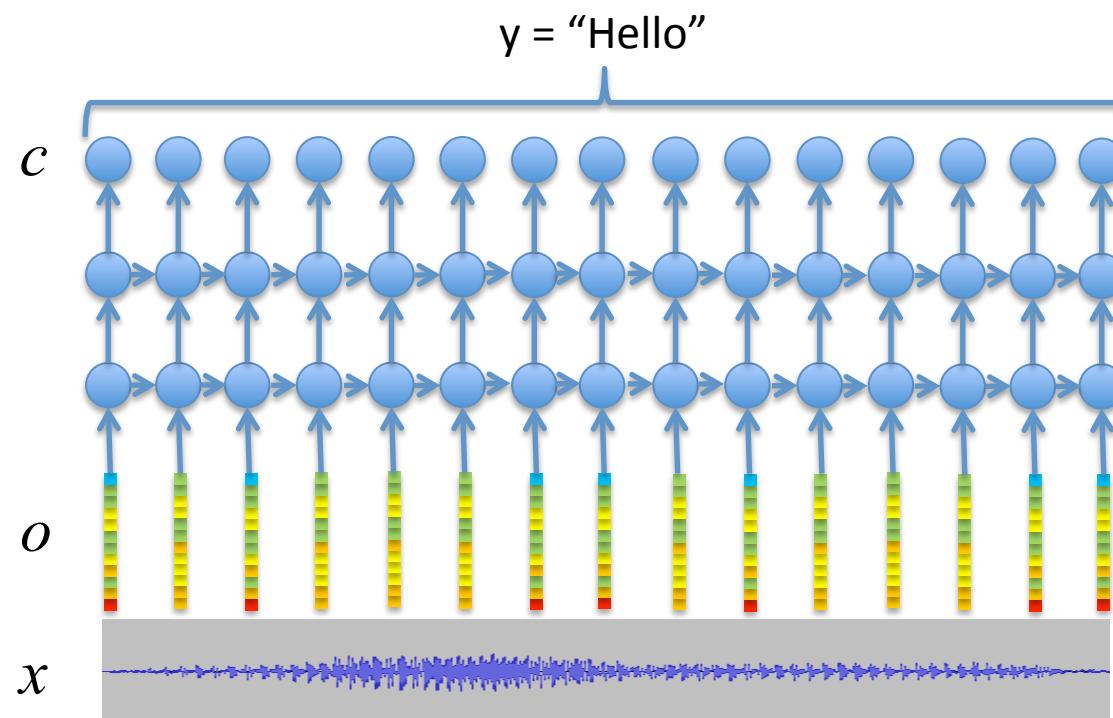
# Spectrogram

- Concatenate frames from adjacent windows to form “spectrogram”.



# Acoustic Model

- Goal: create a neural network (DNN/RNN) from which we can extract transcription,  $y$ .
  - Train from labeled pairs  $(x, y^*)$



# Acoustic model

- Main issue:  $\text{length}(x) \neq \text{length}(y)$ 
  - Don't know how symbols in  $y$  map to frames of audio.
  - Traditionally, try to bootstrap alignment – painful!
- Multiple ways to resolve:
  - Use attention, sequence to sequence models, etc.  
[Chan et al., 2015; Bahdanau et al., 2015]
  - Connectionist Temporal Classification [Graves et al., 2006]

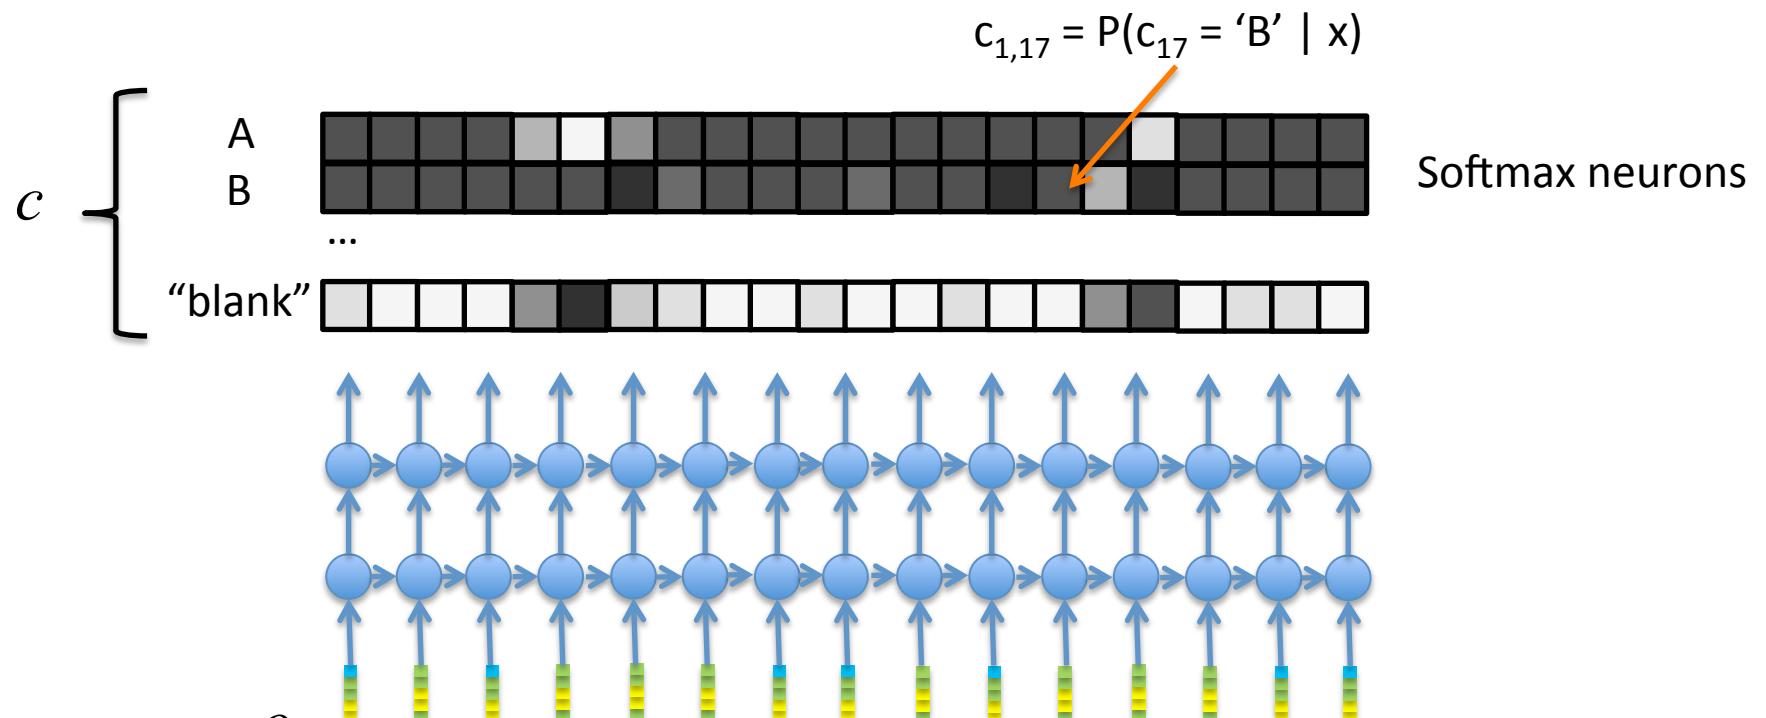
# Connectionist Temporal Classification (CTC)

- Basic idea:
  1. RNN output neurons  $c$  encode distribution over symbols. Note  $\text{length}(c) == \text{length}(x)$ .  
For phoneme-based model:  $c \in \{AA, AE, AX, \dots, ER1, \text{blank}\}$   
For grapheme-based model:  $c \in \{A, B, C, D, \dots, Z, \text{blank}, \text{space}\}$
  2. Define a mapping  $\beta(c) \rightarrow y$ .
  3. Maximize likelihood of  $y^*$  under this model.

# Connectionist Temporal Classification (CTC)

1. RNN output neurons  $c$  encode distribution over symbols. Note  $\text{length}(c) == \text{length}(x)$ .

For grapheme-based model:  $c \in \{A, B, C, D, \dots, Z, \text{blank}, \text{space}\}$



# Connectionist Temporal Classification (CTC)

1. RNN output neurons  $c$  encode distribution over symbols. Note  $\text{length}(c) == \text{length}(x)$ .

For grapheme-based model:  $c \in \{A, B, C, D, \dots, Z, \text{blank}, \text{space}\}$

- Output neurons define distribution over whole character sequences  $c$  assuming independence:

$$P(c|x) \equiv \prod_{i=1}^N P(c_i|x)$$

$$P(c = \text{HHH\_E\_LL\_LO\_\_}|x) = P(c_1 = \text{H}|x)P(c_2 = \text{H}|x) \cdots P(c_{15} = \text{blank}|x)$$

## Connectionist Temporal Classification (CTC)

2. Define a mapping  $\beta(c) \rightarrow y$ .
  - Given a specific character sequence  $c$ , squeeze out duplicates + blanks to yield transcription:

$$y = \beta(c) = \beta(\text{H}\text{H}\text{H\_E\_L}\text{L\_L}\text{O\_}) = \text{"HELLO"}$$

# Connectionist Temporal Classification (CTC)

- Mapping implies a distribution over possible *transcriptions*  $y$ :

$P(c x) = \{$	0.1	HHH_E__LL_L0___	“HELLO” v.1
	0.02	HH_E__LL_L0___	“HELLO” v.2
	0.01	HHH_E__L_L_0H___	“HELL OH”
	0.01	HHH_EE_LL_L_0___	“HELLO” v.3
	...	YY_E__LL_L0_W_	“YELLOW”

$$P(y|x) = \sum_{c:\beta(c)=y} P(c|x)$$

$$P(\text{“HELLO”}) = 0.1 + 0.02 + 0.01 + \dots$$

## Connectionist Temporal Classification (CTC)

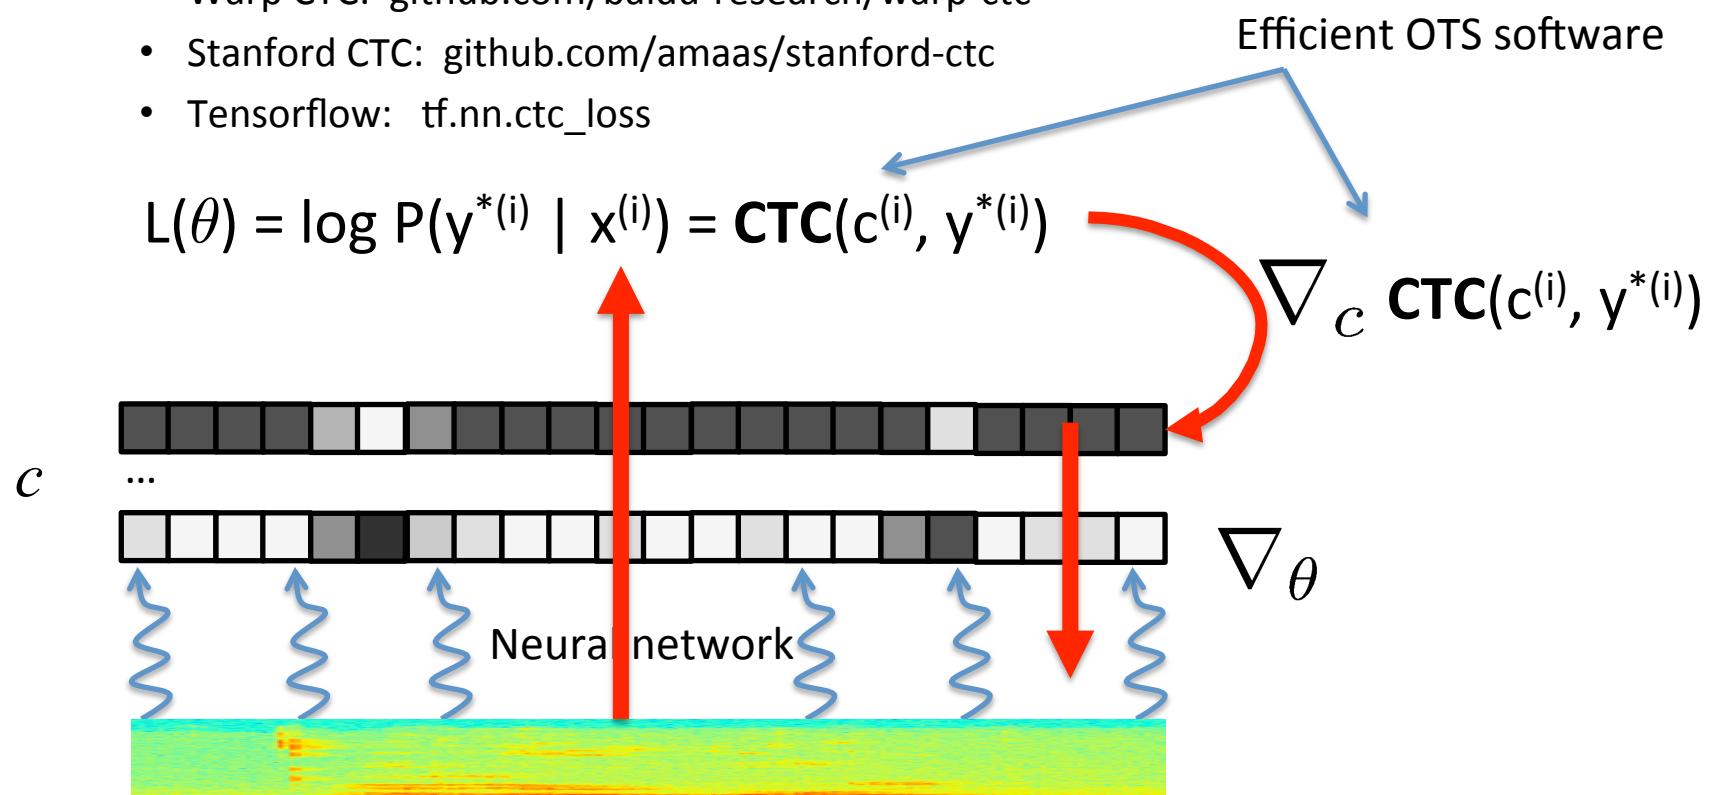
3. Update network parameters  $\theta$  to maximize likelihood of correct label  $y^*$ :

$$\begin{aligned}\theta^* &= \arg \max_{\theta} \sum_i \log P(y^{*(i)} | x^{(i)}) \\ &= \arg \max_{\theta} \sum_i \log \sum_{c: \beta(c)=y^{*(i)}} P(c | x^{(i)})\end{aligned}$$

- [Graves et al., 2006] provides an efficient dynamic programming algorithm to compute the inner summation and its gradient.

# Connectionist Temporal Classification (CTC)

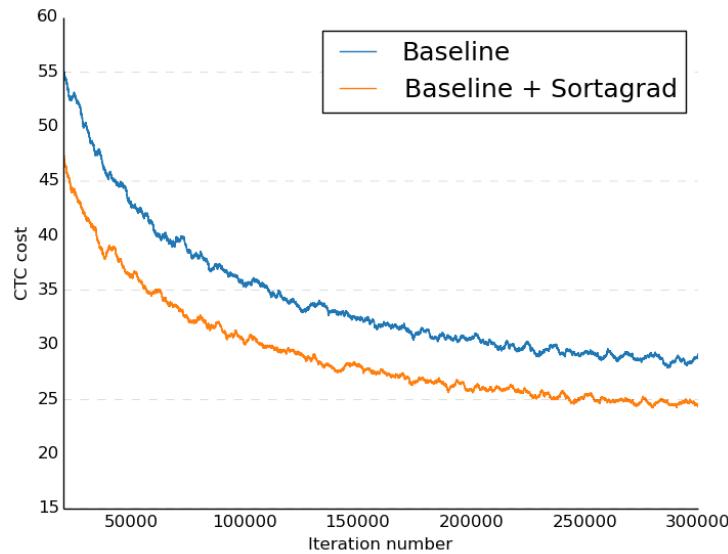
- Use usual gradient descent methods to optimize.  
Tune entire network with backpropagation.
  - Given network outputs, many off-the-shelf packages to compute CTC loss (likelihood) from  $c$  and  $y^*$ , and gradient w.r.t.  $c$ .
    - Warp CTC: [github.com/baidu-research/warp-ctc](https://github.com/baidu-research/warp-ctc)
    - Stanford CTC: [github.com/amaas/stanford-ctc](https://github.com/amaas/stanford-ctc)
    - Tensorflow: `tf.nn.ctc_loss`



# Training tricks

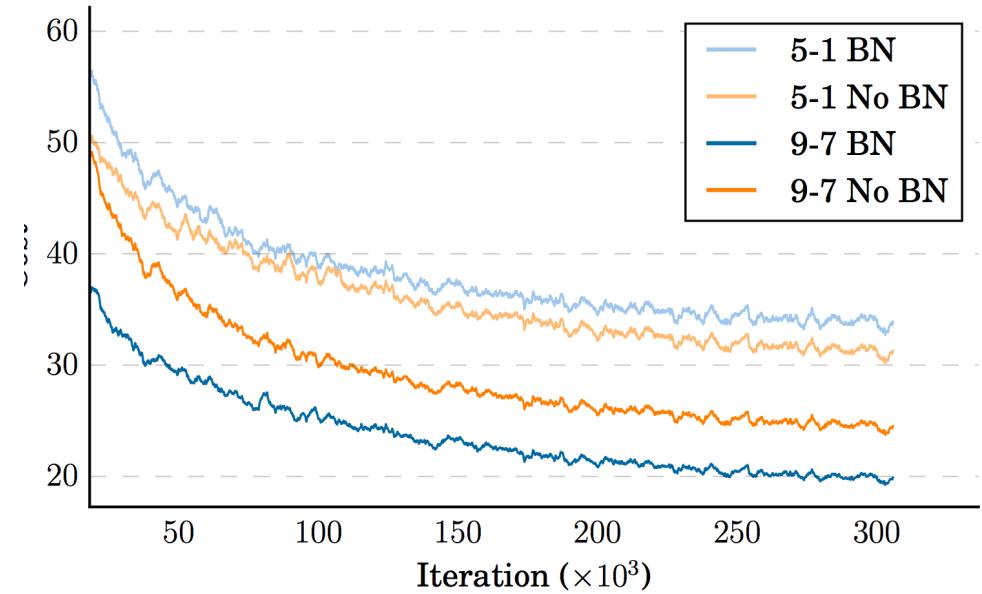
- Getting RNN to train well is tricky.

“SortaGrad”: order utterances by length during first epoch.



See [“Curriculum Learning”, Bengio et al., ICML 2009]

Batch normalization

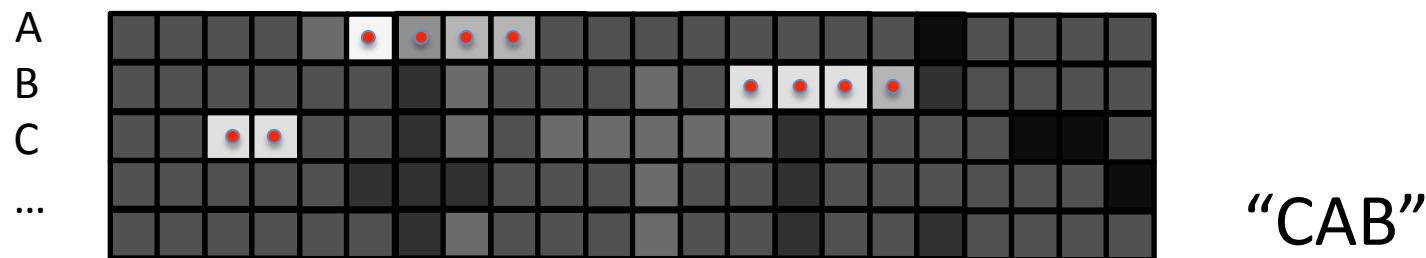


See [Ioffe & Szegedy, 2015]

# Decoding

- Network outputs  $P(c|x)$ . How do we find most likely transcription from  $P(y|x)$ ?
- Simple (approximate) solution:

$$\beta \left( \arg \max_c P(c|x) \right)$$



- Often terrible, but a useful diagnostic to "eyeball" models.

# Example

- Wall Street Journal:
  - <https://catalog.ldc.upenn.edu/Ldc93s6a>
  - Reading WSJ articles.



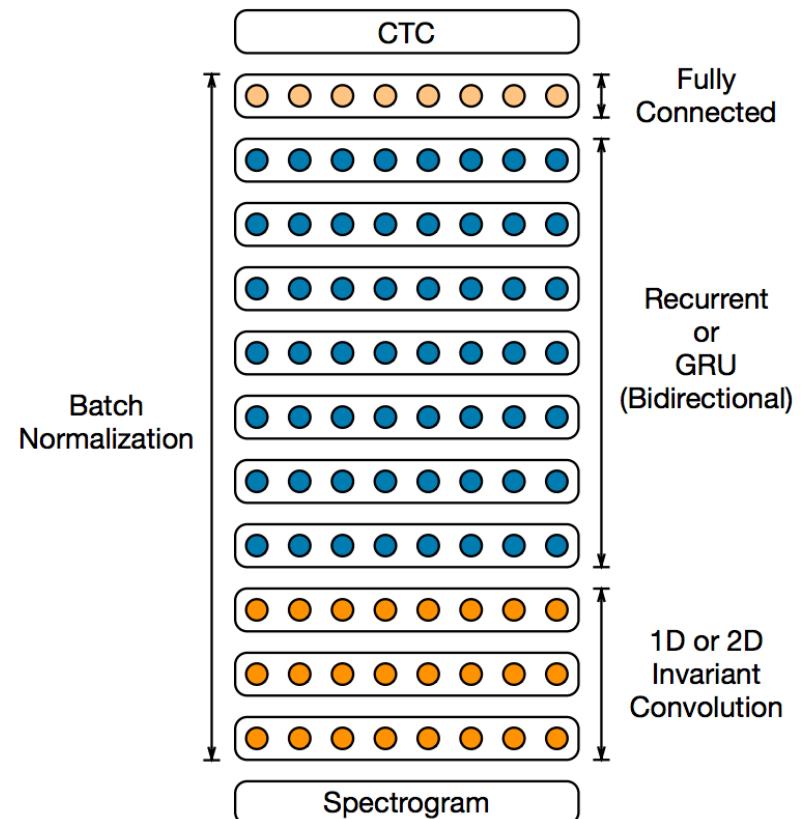
Free alternative: LibriSpeech

- <http://www.openslr.org/12/> [Panayotov et al., ICASSP 2015]
- Read speech from public domain audiobooks.

# Example

- RNN to predict graphemes (26 characters + space + blank):
  - Spectrograms as input.
  - 1 layer of convolutional filters.
  - 3 layers of Gated Recurrent Units.
    - 1000 neurons per layer.
  - 1 full-connected layer to predict  $c$ .
  - Batch normalization
- CTC loss function (Warp-CTC)
- Train with SGD+Nesterov momentum.

Typical model family:

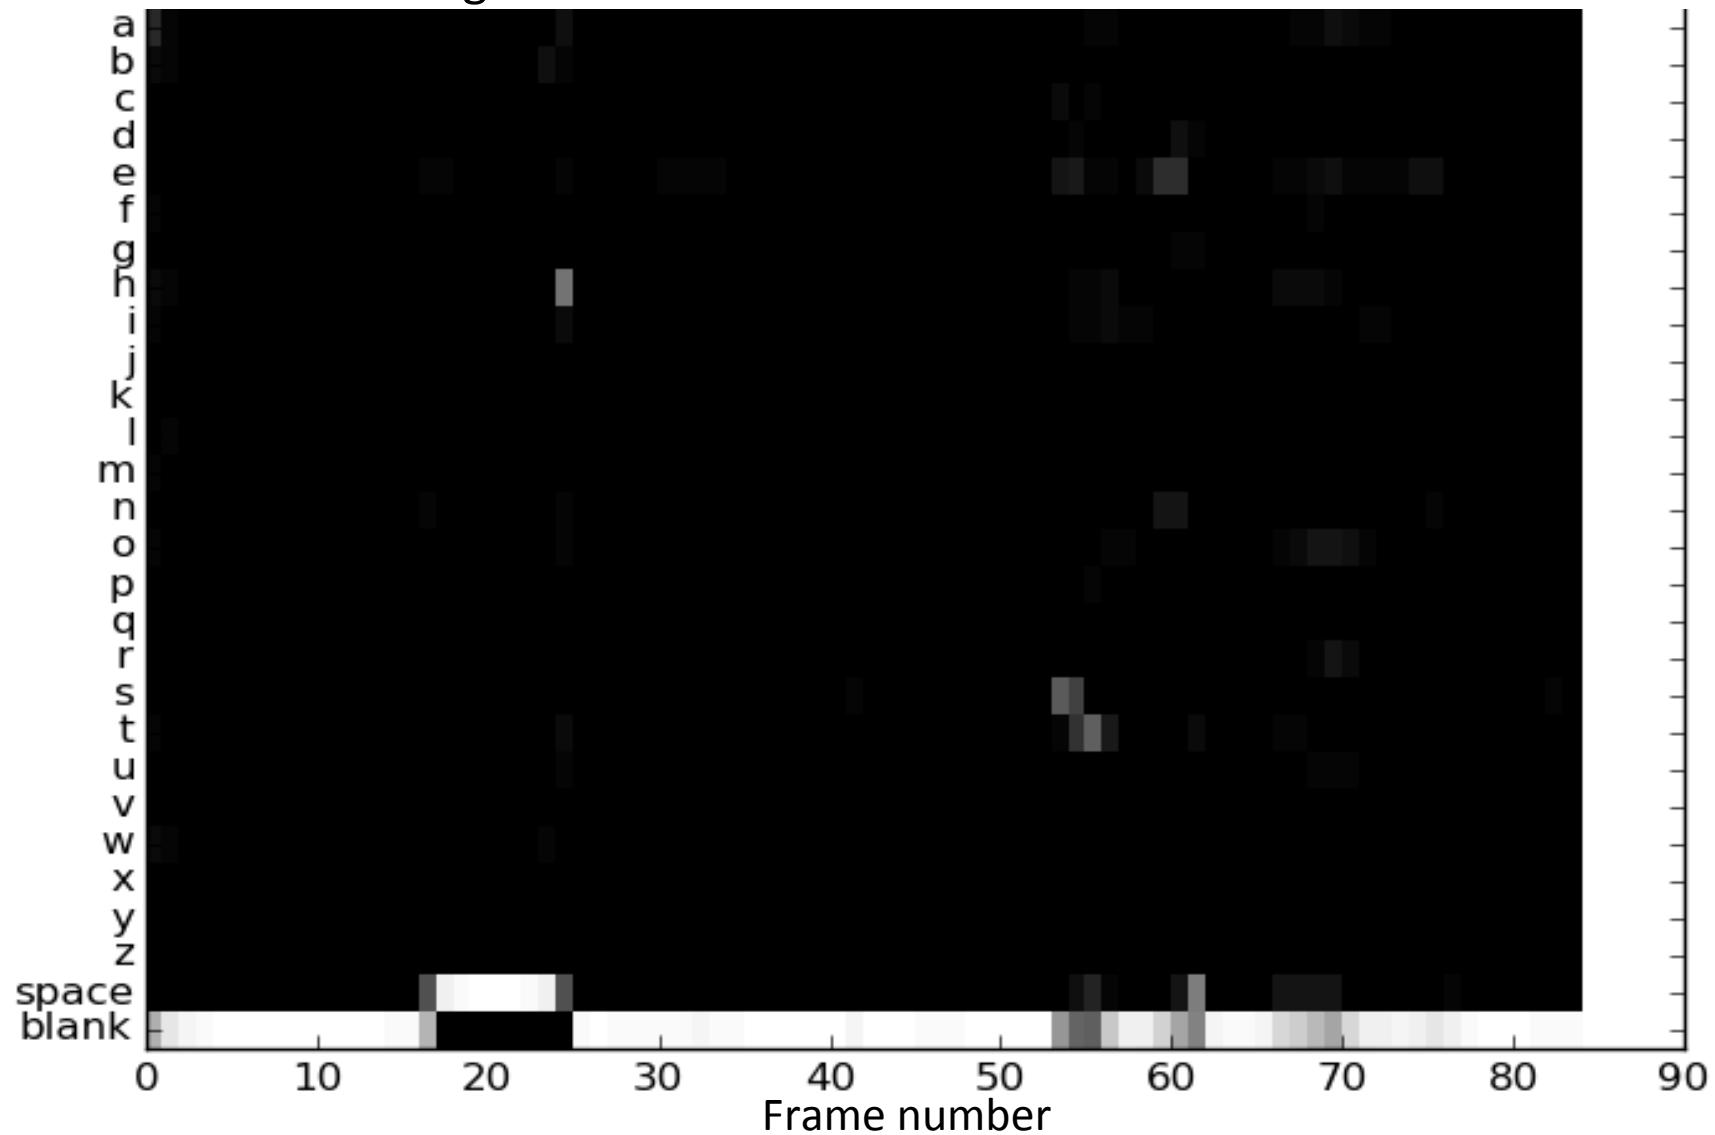


# Example

- What happens inside?

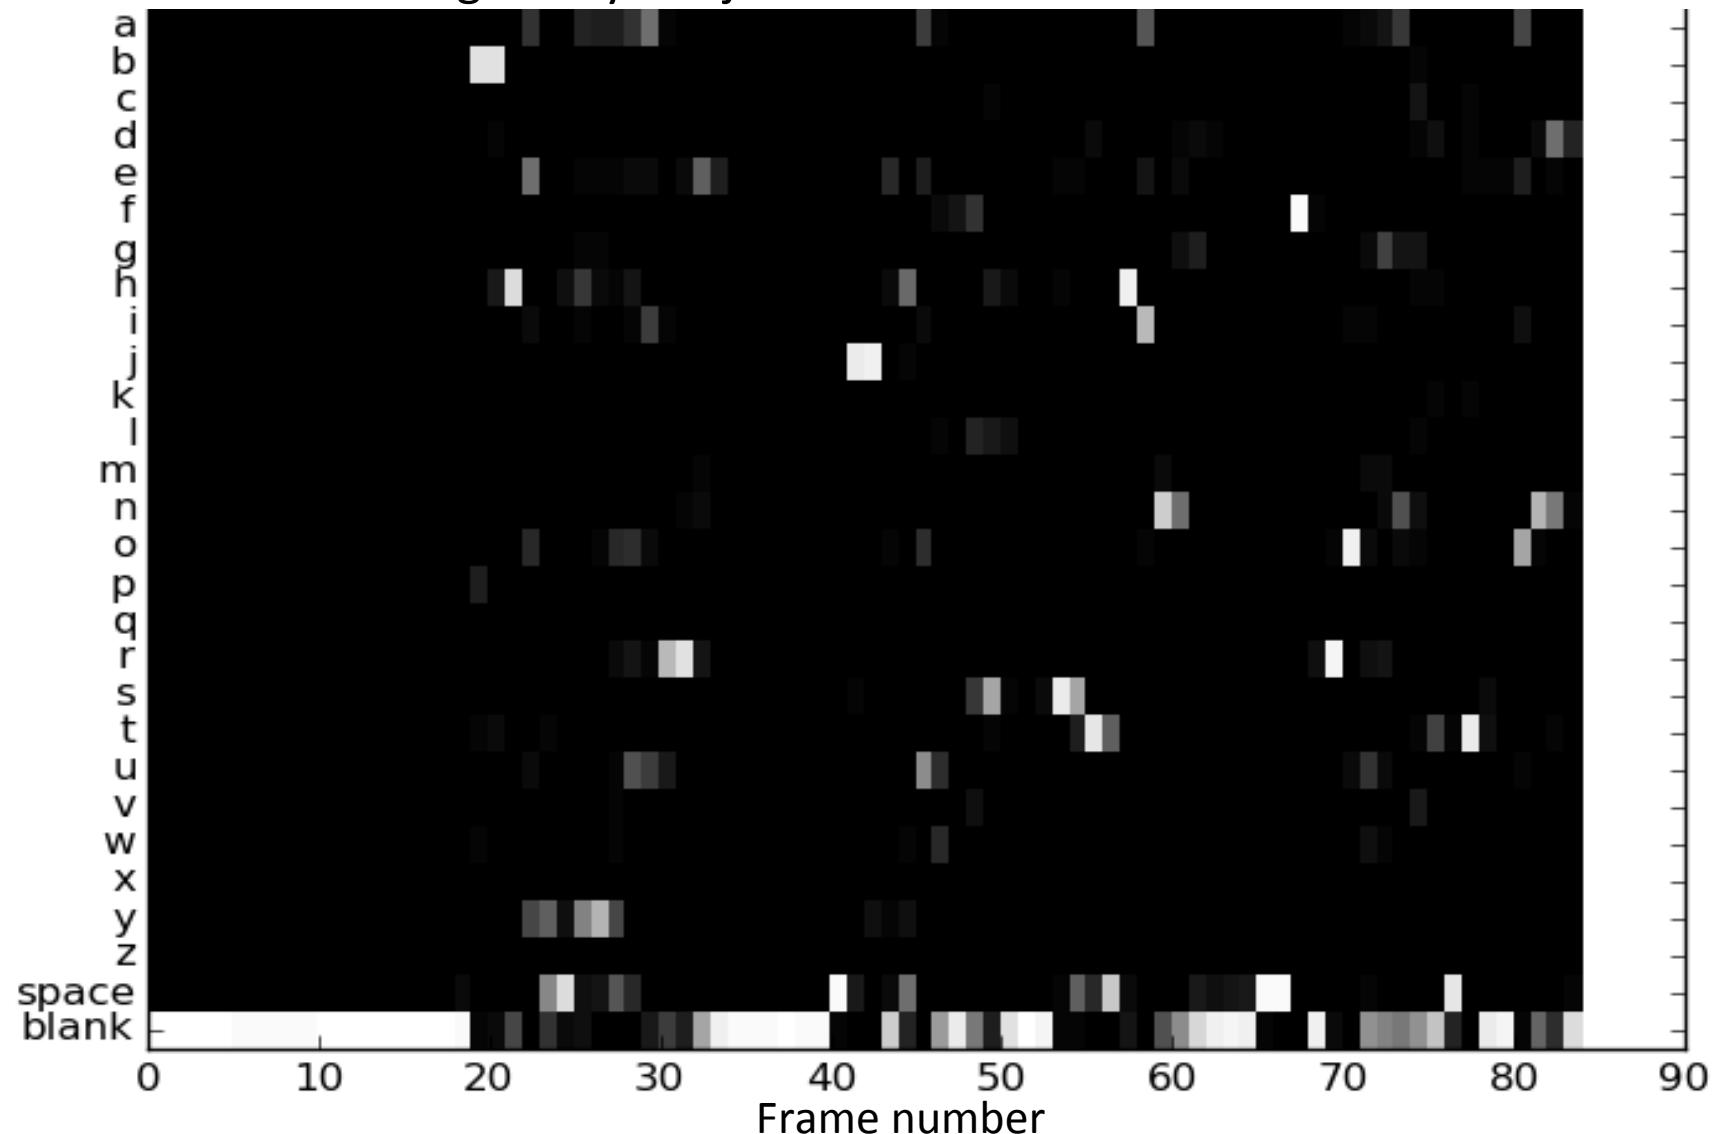
Network outputs ( $c$ ) at Iteration 300  
(Thresholding / contrast added for clarity.)

Max decoding: h



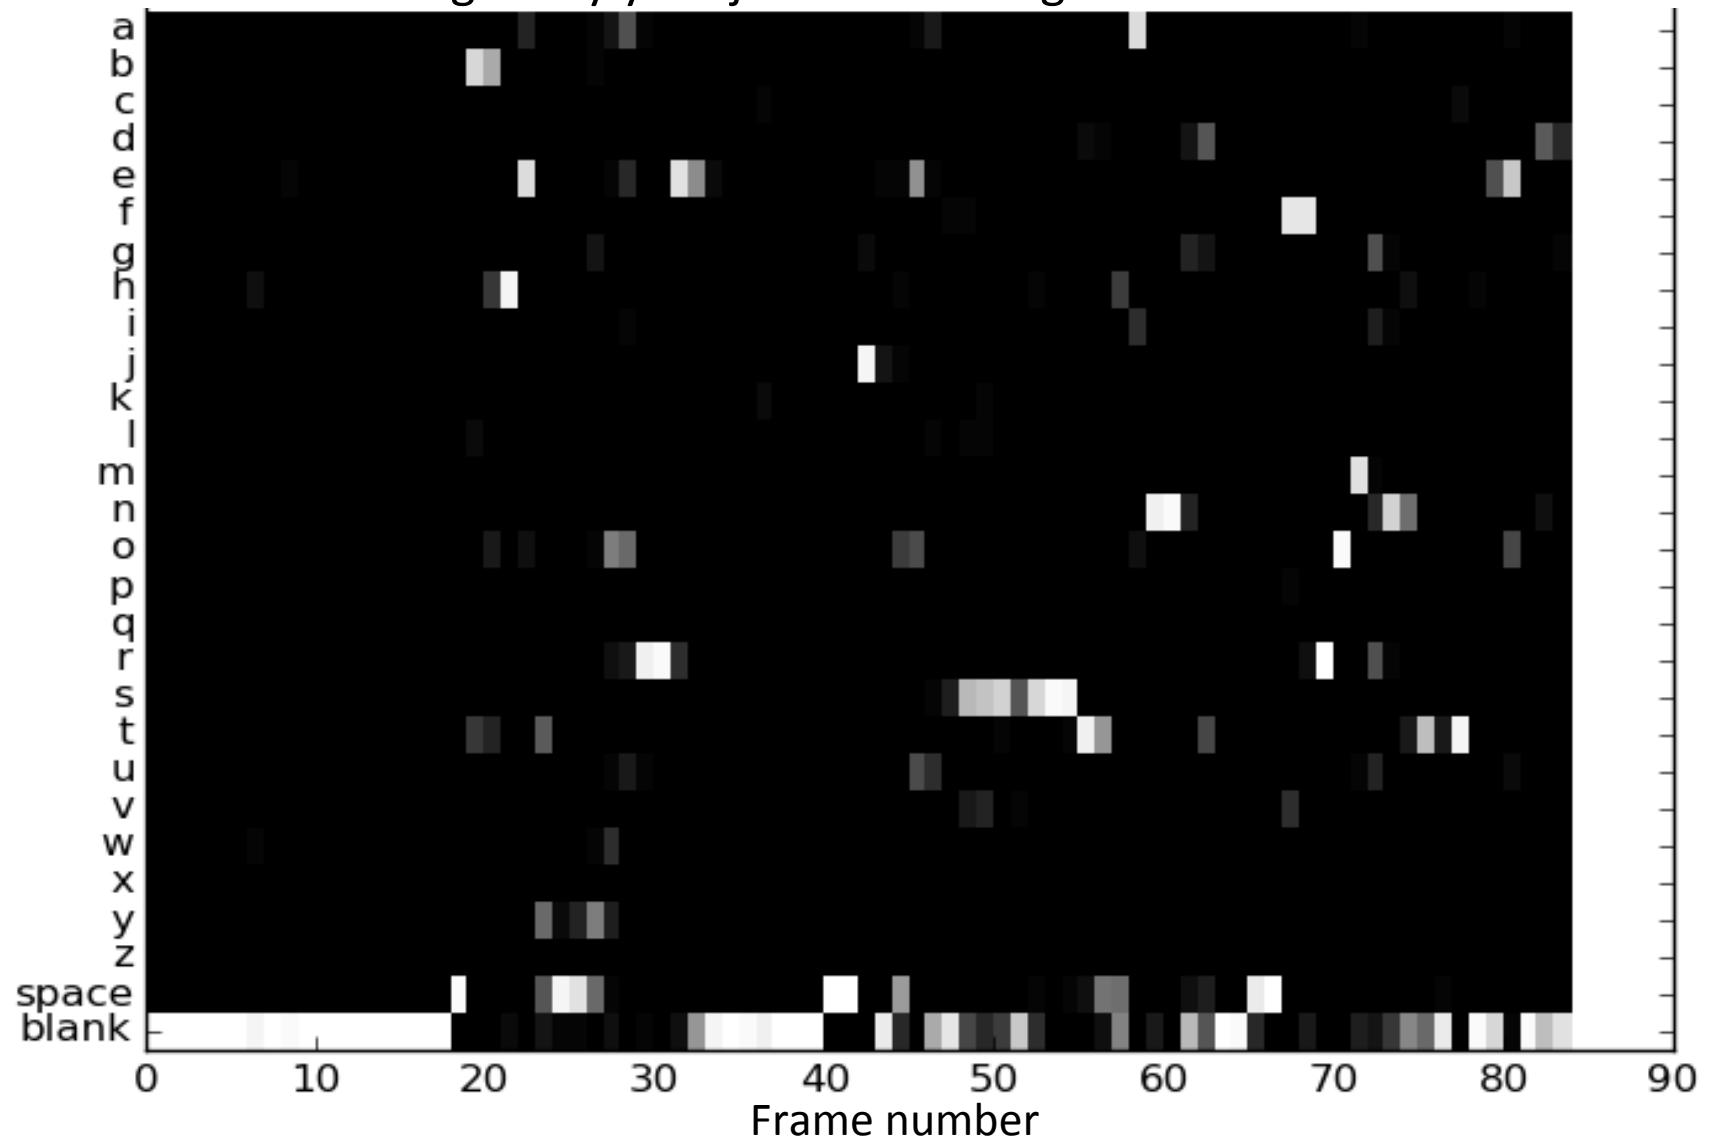
# Network outputs ( $c$ ) at Iteration 1500 (Thresholding / contrast added for clarity.)

Max decoding: bhe y uar j usst hin fro ton



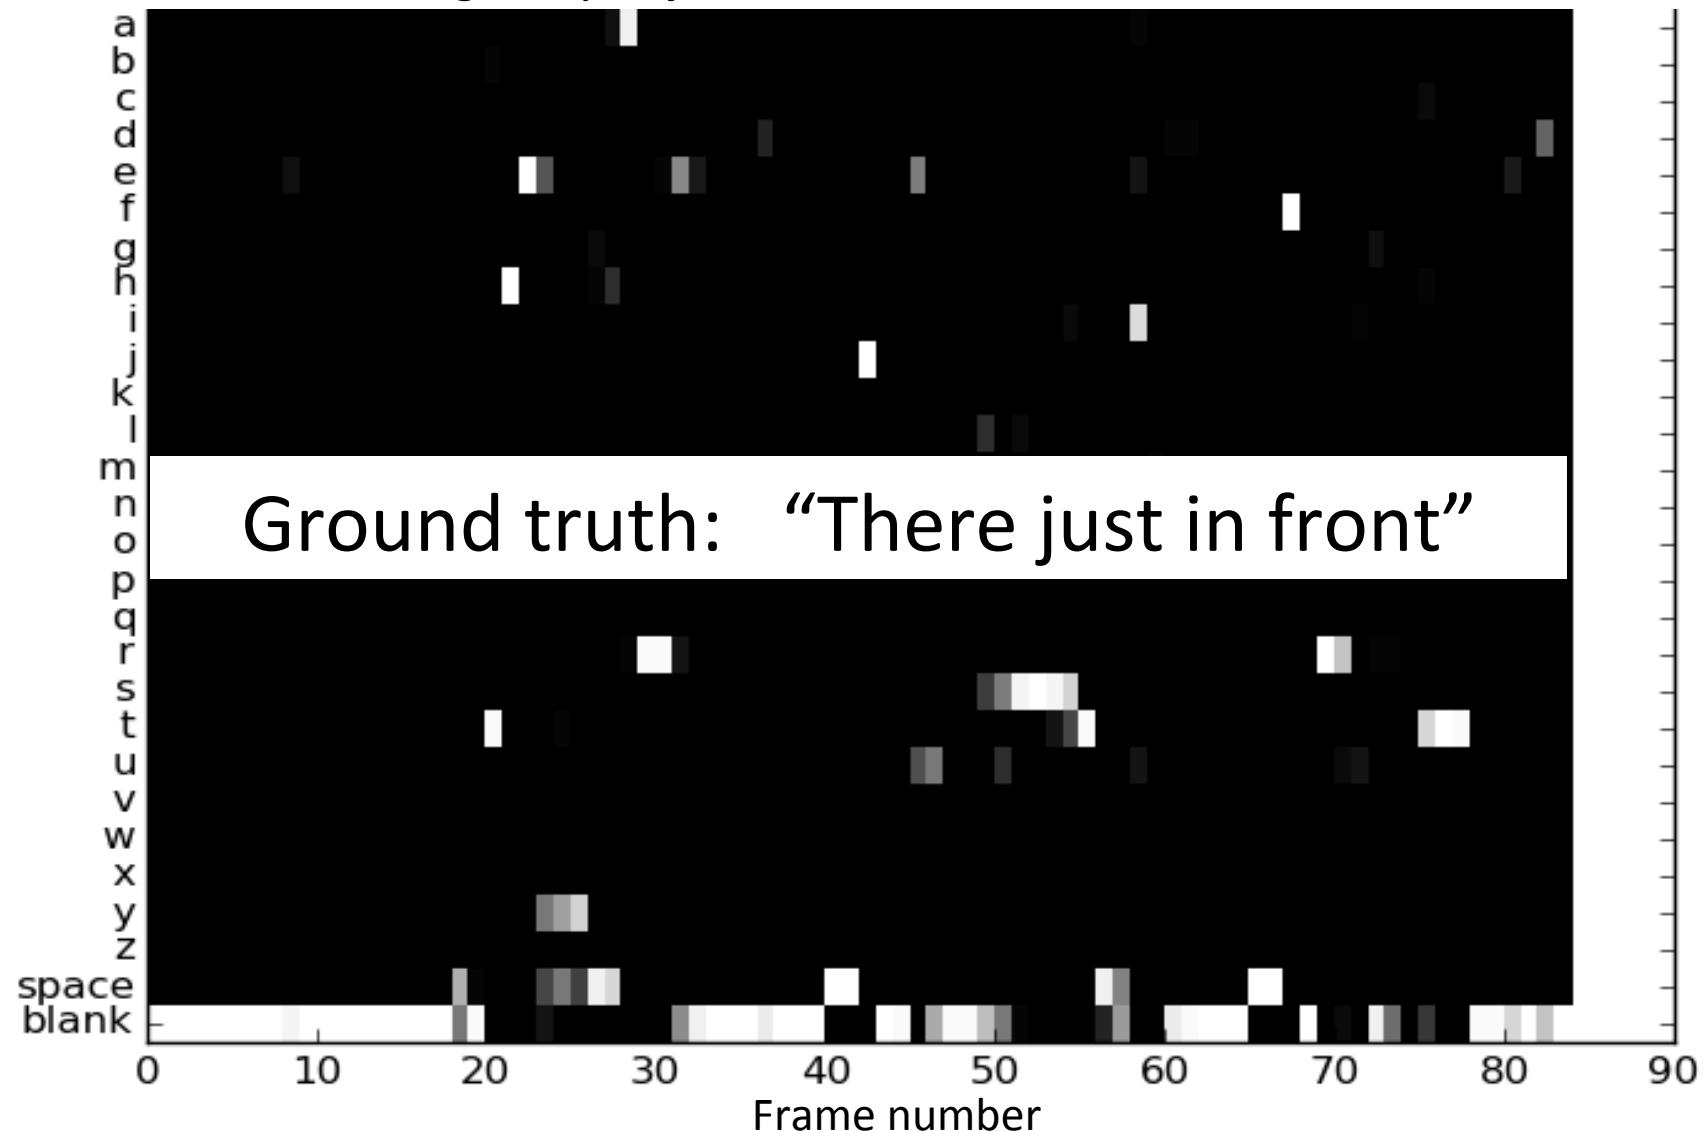
# Network outputs ( $c$ ) at Iteration 2500 (Thresholding / contrast added for clarity.)

Max decoding: bhey yore j esstand fromgntte



Network outputs ( $c$ ) at Iteration 5500  
(Thresholding / contrast added for clarity.)

Max decoding: they ar jest in front



# Max Decoding

- Examples:



**Max Decoding**

“put pore lottle thank and sr crits sinpt the atting to them  
having been turned ef the wal al thes years con”

**True Label**

“the poor little things cried cynthia think of them  
having been turned to the wall all these years”



**Max Decoding**

“that is true baddel gre”

**True Label**

“that is true badauderie”

# Decoding

- Network outputs  $P(c|x)$ . How do we find most likely transcription from  $P(y|x)$ ?
- No efficient solution in general. Resort to search!
  - See [Graves et al., 2006] for prefix decoding strategy.

# Language models

- Even with better decoding, CTC model tends to make spelling + linguistic errors. E.g.:

RNN output

---

what is the weather like in bostin right now  
prime miniter nerenr modi  
arther n tickets for the game

From Hannun et al., 2014.

- $P(y|x)$  modeled directly from audio.
  - But not enough audio data to learn complicated spelling and grammatical structure.
  - Only supports small vocabulary.
  - For grapheme models: “Tchaikovsky” problem.

# Language models

- Two solutions
  - Fuse acoustic model with language model:  $P(y)$
  - Incorporate linguistic data:
    - Predict phonemes + pronunciation lexicon + LM.
- Possible to train language model from massive *text corpora*.
  - Learn spelling + grammar
  - Greatly expand vocabulary
  - Elevate likely cases (“Tchaikovsky concerto”) over unlikely cases (“Try cough ski concerto”).

# Language models

- Standard approach: n-gram models
  - Simple n-gram models are common, well supported.
    - KenLM: [kheafield.com/code/kenlm/](http://kheafield.com/code/kenlm/)
  - Train easily from huge corpora.
  - Quickly update to follow trends in traffic.
  - Fast lookups inside decoding algorithms.

# Decoding with LMs

- Given a word-based LM of form  $P(w_{t+1} | w_{1:t})$ , Hannun et al., 2014 optimize:

$$\arg \max_w P(w|x)P(w)^\alpha [\text{length}(w)]^\beta$$

$P(w|x) = P(y|x)$  for characters that make up  $w$ .

$\alpha$  and  $\beta$  are tunable parameters to govern weight of LM and a bonus/penalty for each word.

# Decoding with LMs

- Basic strategy: beam search to maximize

$$\arg \max_w P(w|x)P(w)^\alpha [\text{length}(w)]^\beta$$

Start with set of candidate transcript prefixes,  $A = \{\}$ .

For  $t = 1..T$ :

For each candidate in  $A$ , consider:

1. Add blank; don't change prefix; update probability using AM.
2. Add space to prefix; update probability using LM.
3. Add a character to prefix; update probability using AM.

Add new candidates with updated probabilities to  $A_{\text{new}}$ .

$A := K$  most probable prefixes in  $A_{\text{new}}$ .

# Decoding with LMs: Examples

RNN output	Decoded Transcription
what is the weather like in bostin right now prime miniter nerenr modi arther n tickets for the game	what is the weather like in boston right now prime minister narendra modi are there any tickets for the game

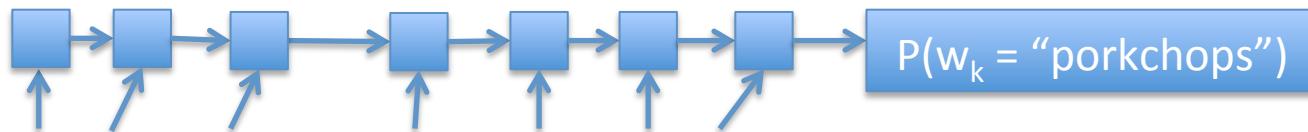
From Hannun et al., 2014.

# Rescoring

- Another place to plug in DL algorithms:  
Systems usually produce N-best list.  
Use fancier models to “rescore” this list.

# Rescoring with Neural LM

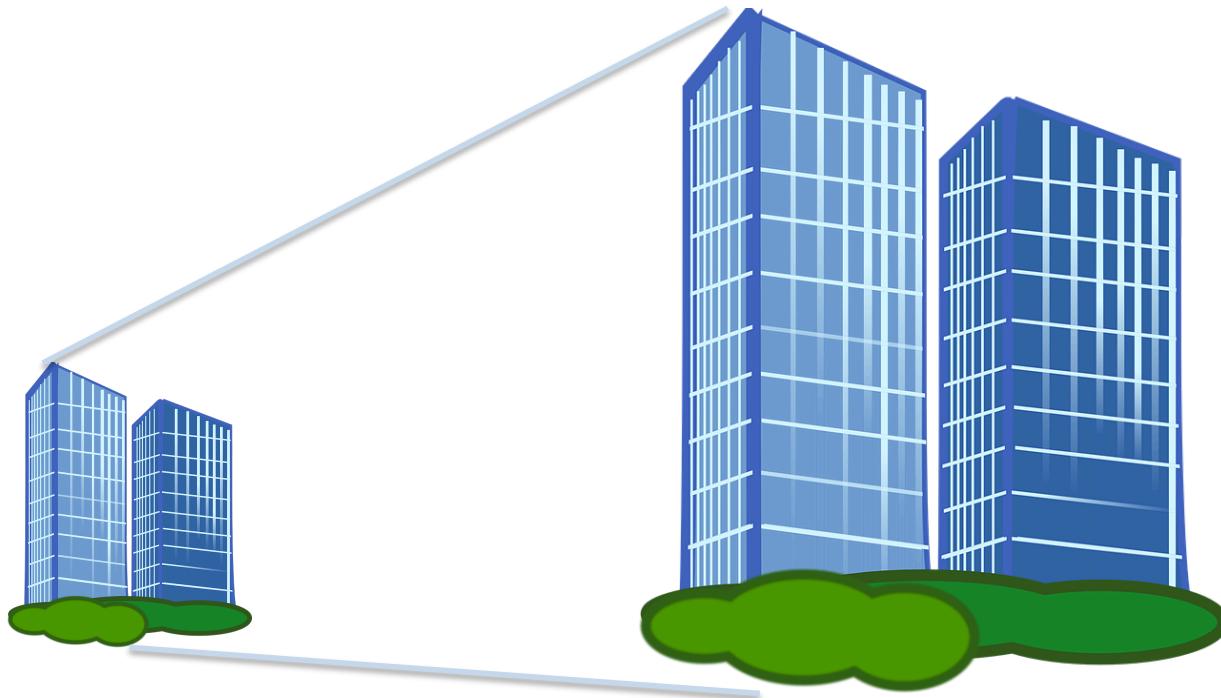
- Example: train neural language model and rescore word transcriptions.
  - Cheap to evaluate  $P(w_k | w_{k-1}, w_{k-2}, \dots, w_1)$  NLM on many sentences.
  - In practice, often combine with N-gram trained from big corpora.



1. (-25.45) I'm a connoisseur looking for wine and porkchops. -24.45
2. (-26.32) I'm a connoisseur looking for wine and port shops. -23.45
3. ...
4. ...
5. ...

# **SCALING UP**

# Scale model



- Two main components to scale from “tutorial” to state-of-art accuracy:
  - Data
  - Computing power

# Data

- Transcribing speech data isn't cheap, but not prohibitive:
  - Roughly 50¢ to \$1 per minute.
- Typical speech benchmarks offer 100s to few 1000s of hours.
  - LibriSpeech (audiobooks)
  - LDC corpora (Fisher, Switchboard, WSJ) (\$\$)
  - VoxForge

# Types of speech data

- Application matters
  - We want to find data that matches our goals.

Styles of speech	Issues	Applications
Read Conversational Spontaneous Command/control	Disfluency / stuttering Noise Mic quality / #channels Far field Reverb / echo Lombard effect Speaker accents	Dictation Meeting transcription Call centers Device control Mobile texting Home / IoT / Cars

# Read speech

- Reading is inexpensive way to get more data.  
< \$10/hour depending on source
- Disadvantages:
  - Misses inflection/conversational tone
  - Lombard effect
  - Speaker variety sometimes a limitation.

# Read speech

- Some tricks we've tried to address these:
  - Elicit “voice acting” by using movie scripts:



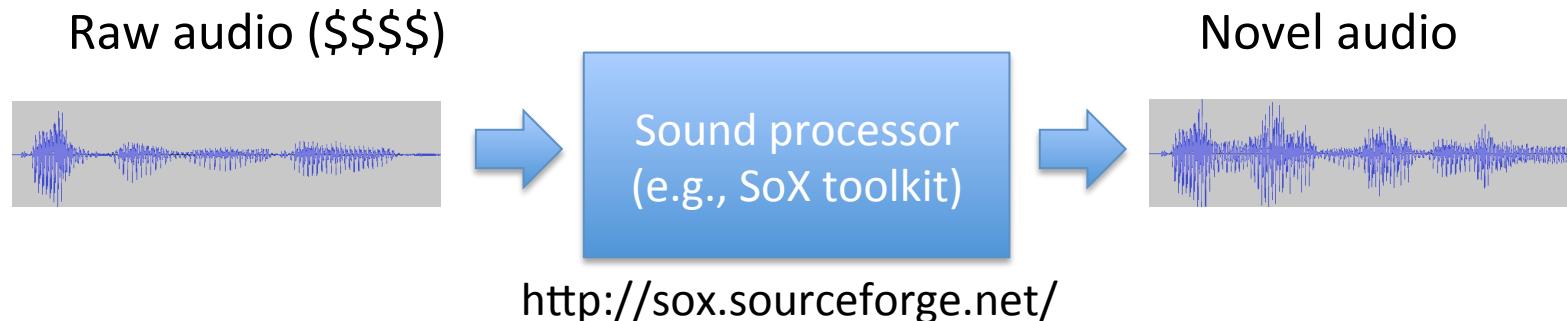
- Elicit Lombard effects by playing loud noise (sometimes via headphones):



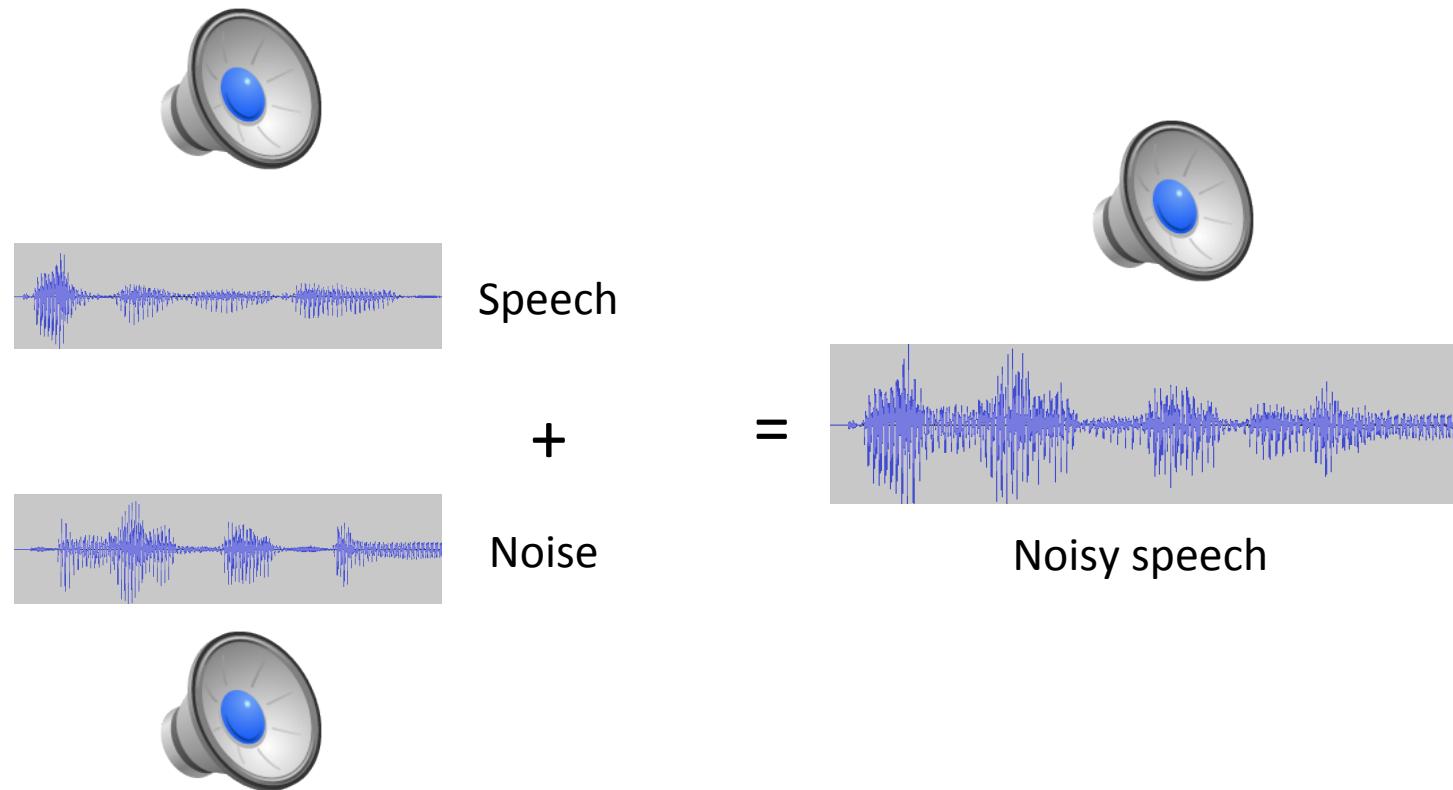
[Sanjeev Satheesh]

# Augmentation

- Many forms of distortion that model should be robust to:
  - Reverb, noise, far field effects, echo, compression artifacts, changes in tempo



# Example: additive noise



DeepSpeech 2: 10K hours of raw audio -> 100K hours of novel training data

# Example: tempo



WSJ example



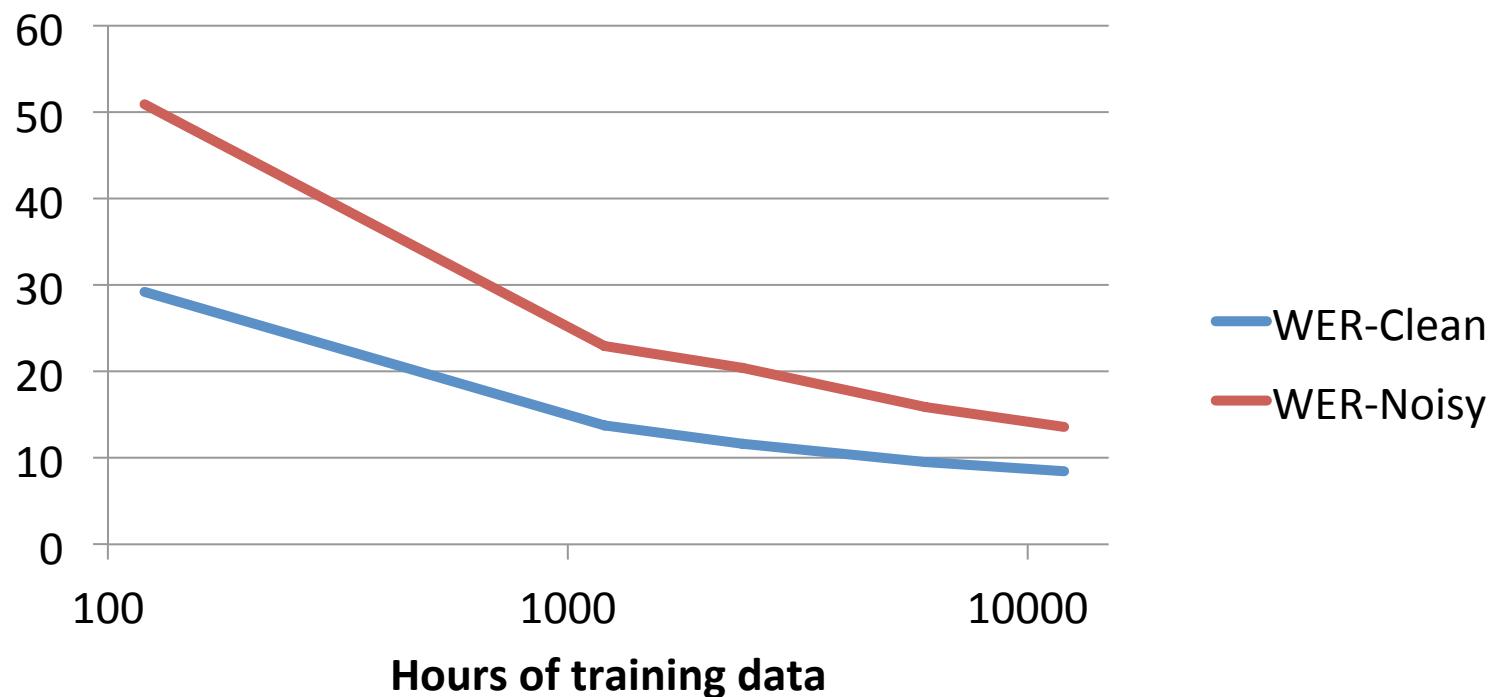
Faster WSJ reader w/ reverb

```
sox <infile> <outfile> tempo 1.3 reverb
```

- In general: easier to engineer data pipeline than to engineer recognition pipeline.

# Results: DeepSpeech 2

- Steady fall in error rates with new raw data.



(Assuming we have a big enough model!)

# Computation

- How big is 1 experiment?

At least:

$$(\# \text{ connections}) \cdot (\# \text{ frames}) \cdot (\# \text{ utterances}) \cdot \\ (\# \text{ epochs}) \cdot 3 \cdot 2 \text{ FLOPs}$$

E.g.: for DS2 with 10K hours of data:

$$100\text{e}6 \cdot 100 \cdot 10\text{e}6 \cdot 20 \cdot 3 \cdot 2 = 1.2\text{e}19 \text{ FLOPs}$$

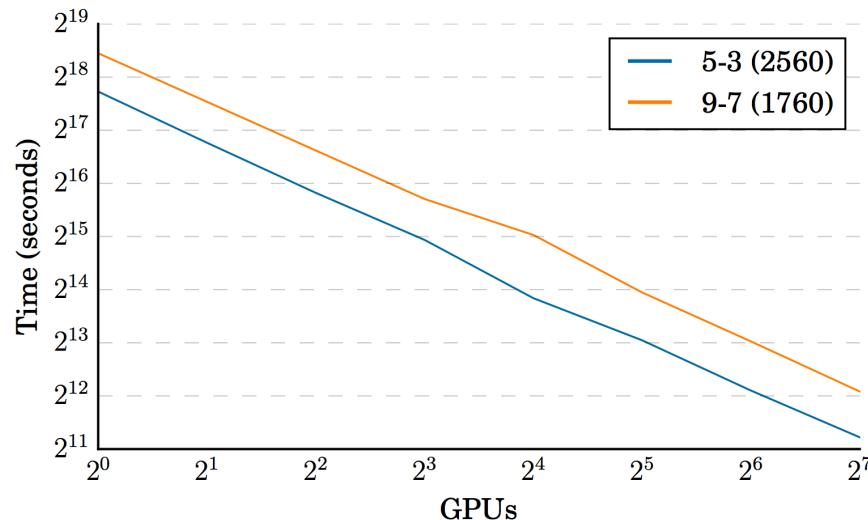
~30 days with well-optimized code on Titan X.

# Computation

- Easy: use more GPUs with data parallelism.
  - Minibatches up to 1024 seem useful.
  - Would like  $\geq 64$  utterances per GPU.
- Current servers support ~8 Titans.
  - Will get you < 1 week training time.

# Computation: multi-node

- Many ways to use more GPUs.
  - At Baidu, use synch. SGD:

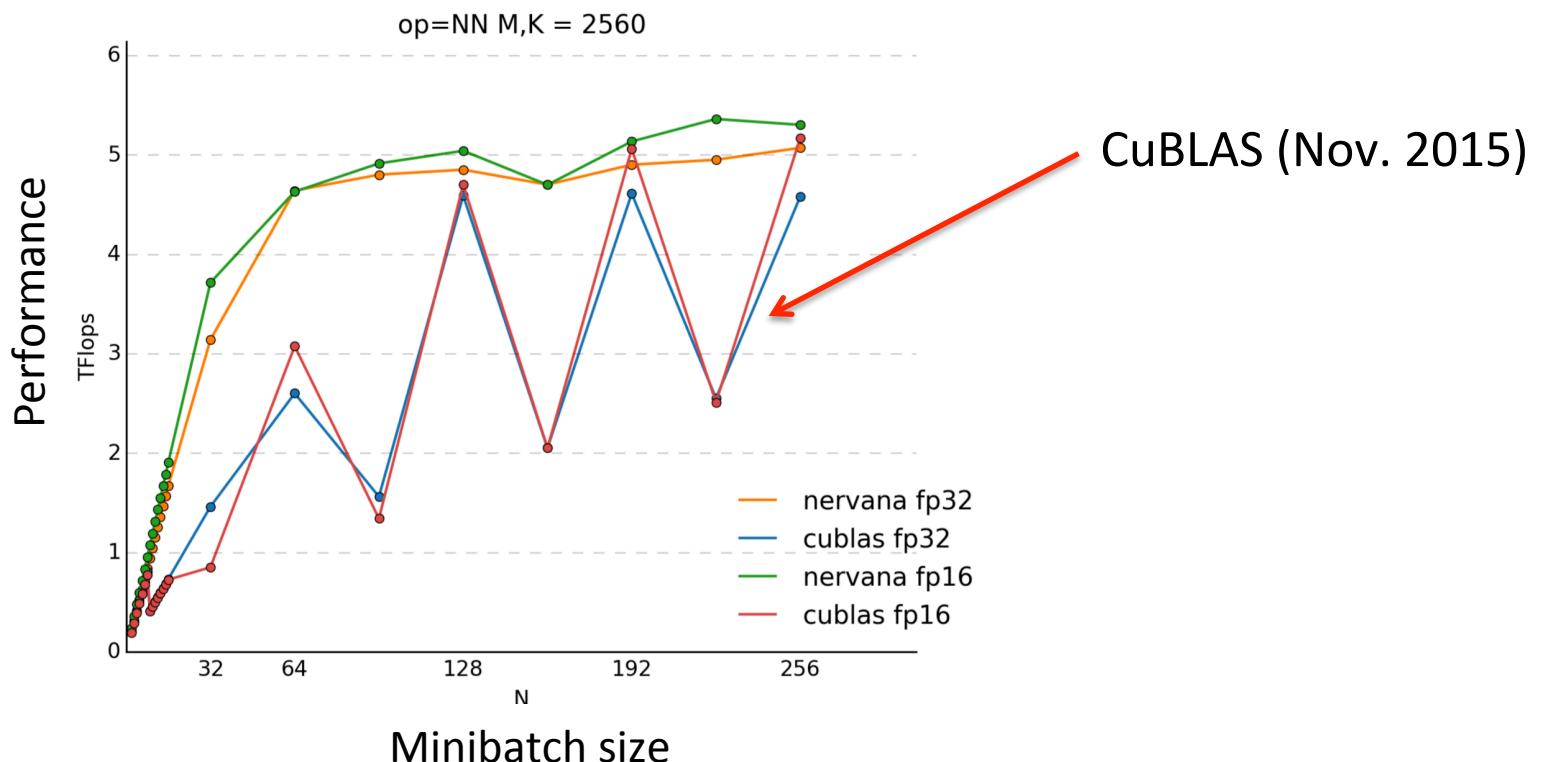


Needed to optimize OpenMPI to achieve efficiency.

- Other solutions:
  - Async SGD [Dean et al., NIPS 2012]
  - Synch SGD w/ backup workers [Chen et al., ICLR 2016]

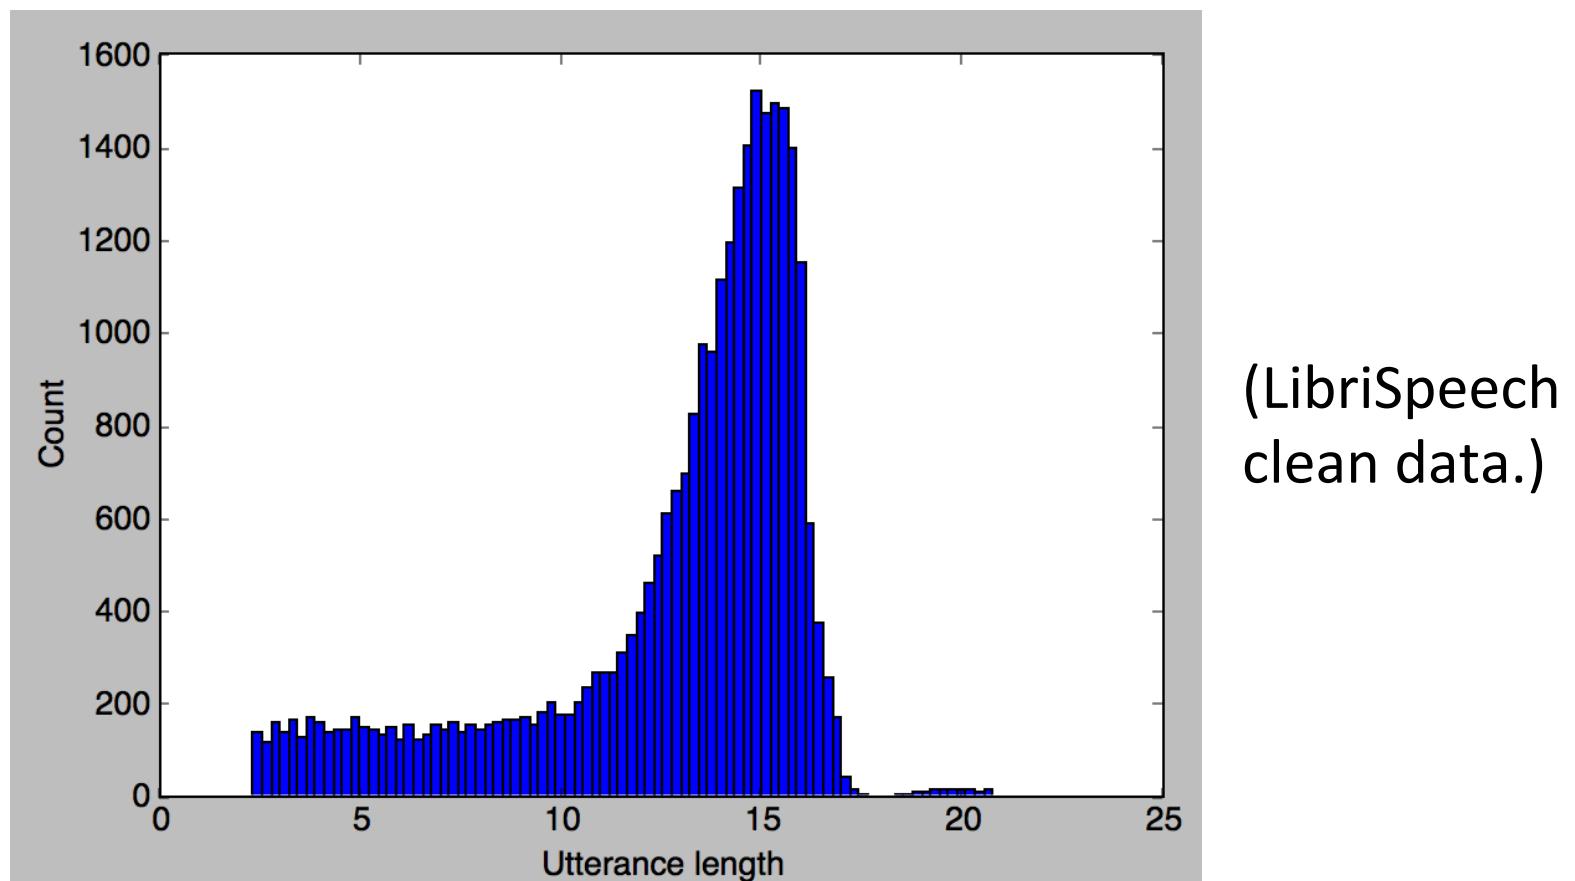
# Computation

- Need *optimized* single-GPU code. But a lot of off-the-shelf code has inefficiencies.
- E.g., Watch for bad GEMM sizes:



# Computation

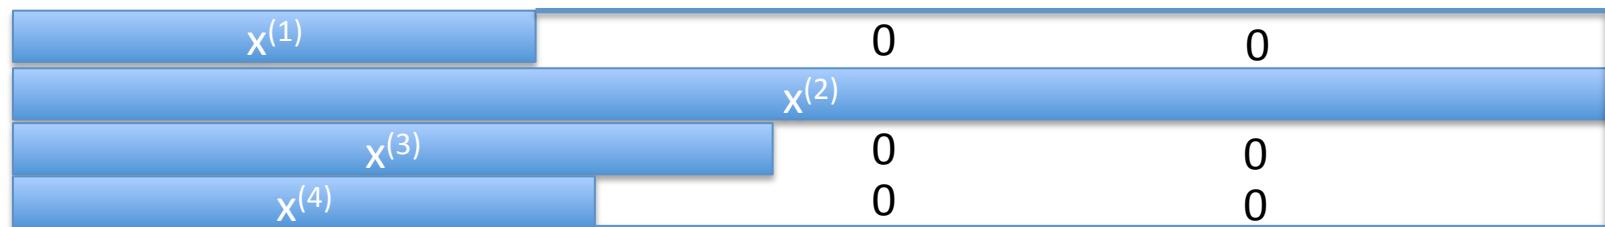
- Try to keep similar-length utterances together.



# Computation

- Try to keep similar-length utterances together.

Bad minibatch:

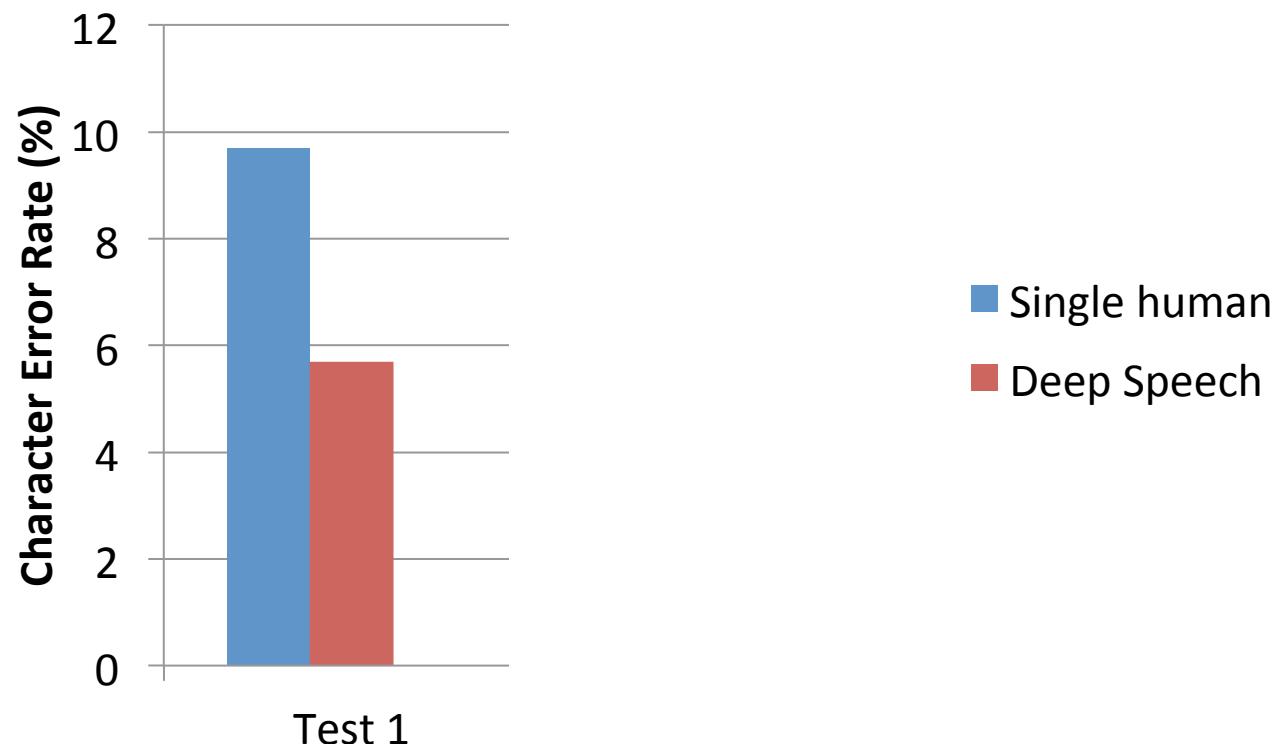


Good minibatch:



# Results

- Scaled up models in Mandarin:



Reality check

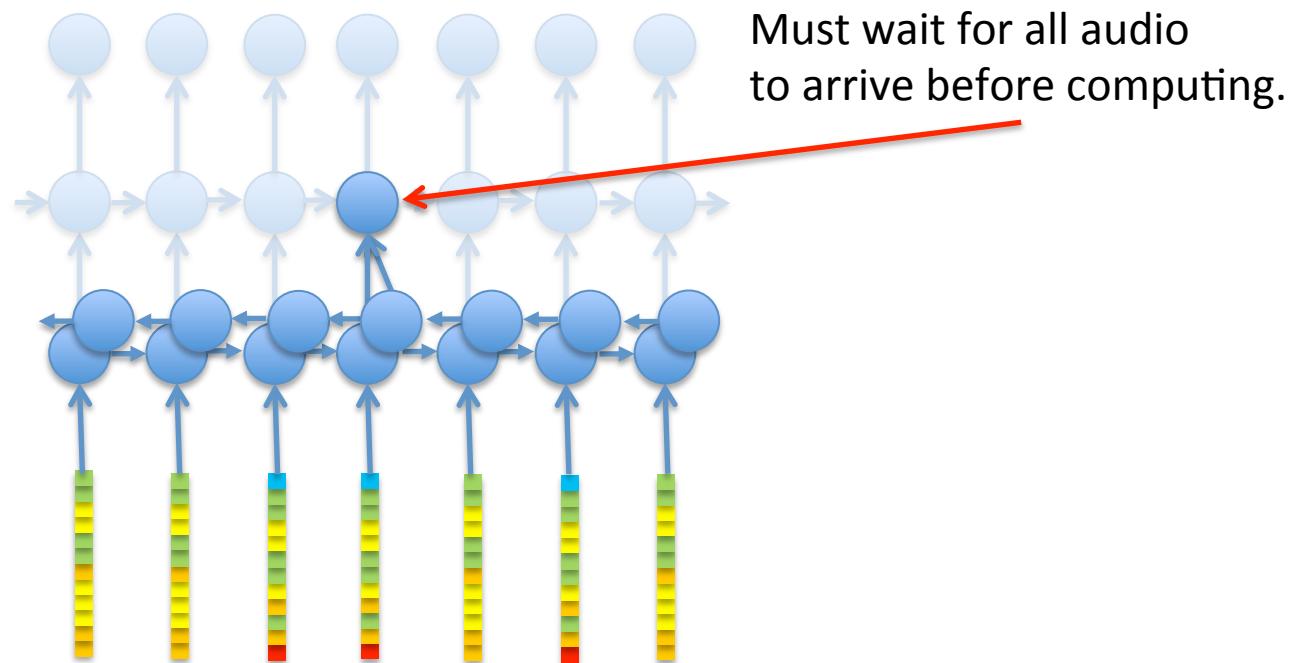
# **PRODUCTION**

# Production speech

- So far:
  - Train acoustic + language models.
  - Scale them up.
- But how to serve users?
  - Accuracy is only one measure of performance.
  - Users care about latency.
  - Need to serve economically.

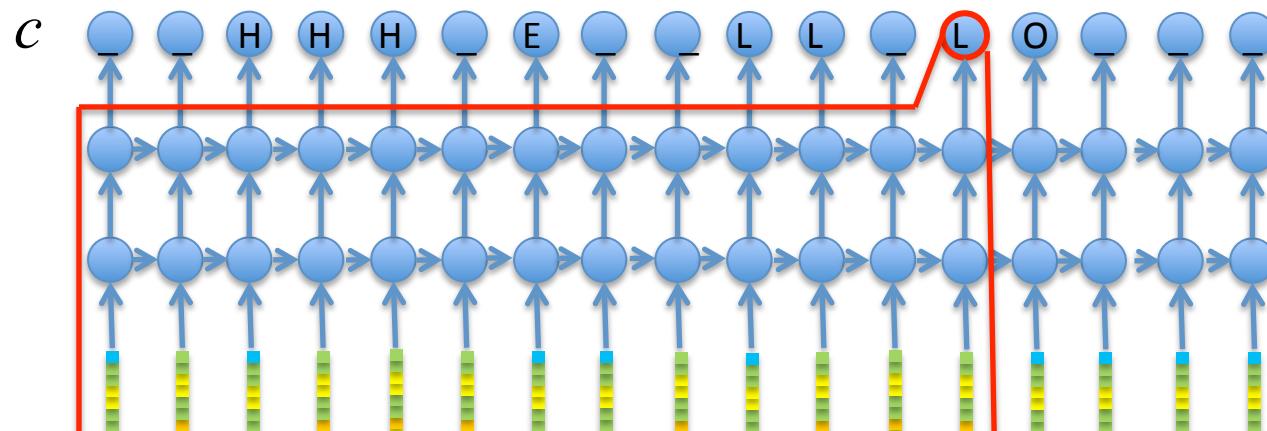
# Latency

- Many acoustic model structures hard to serve in practice.
  - E.g., Bi-directional RNNs.



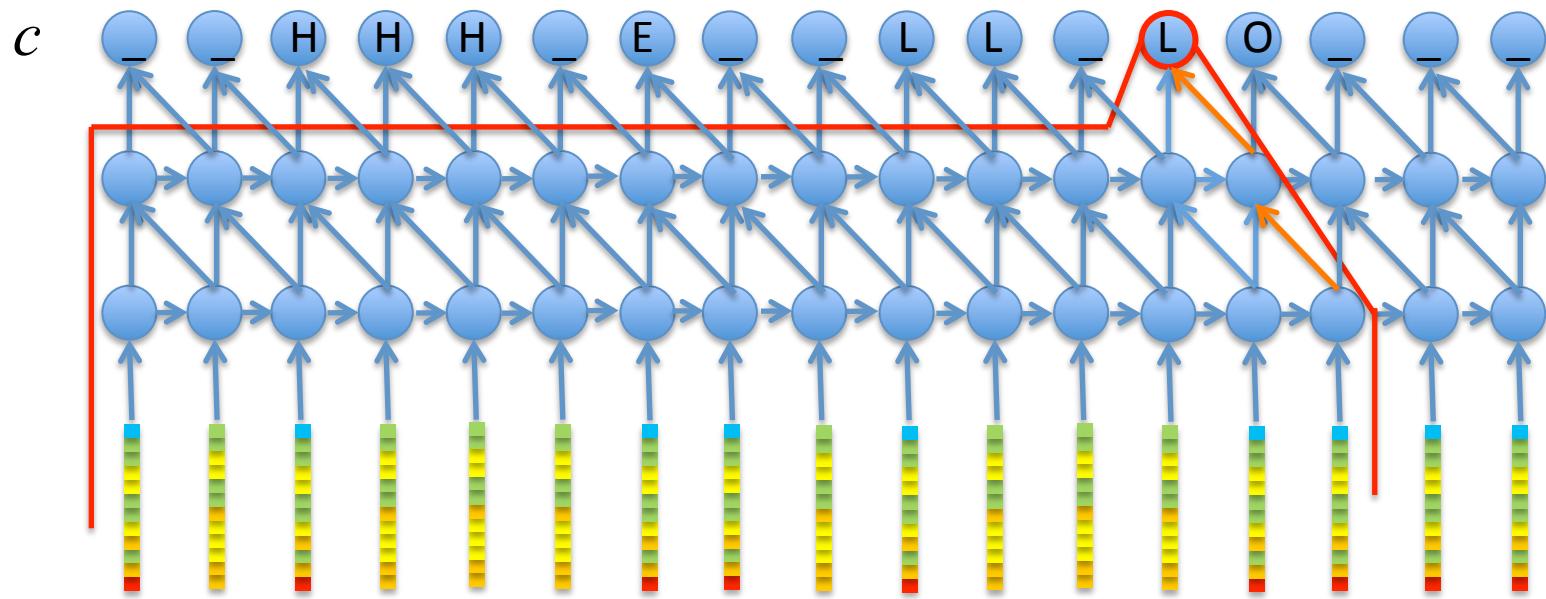
# Latency

- Use forward-only RNNs. But:
  - Usually hurts accuracy. Why? Context?
  - CTC could learn to delay output on its own in order to improve accuracy.
    - In practice, tends to align transcription closely.
    - This is especially problematic for English letters (spelling!)



# Latency

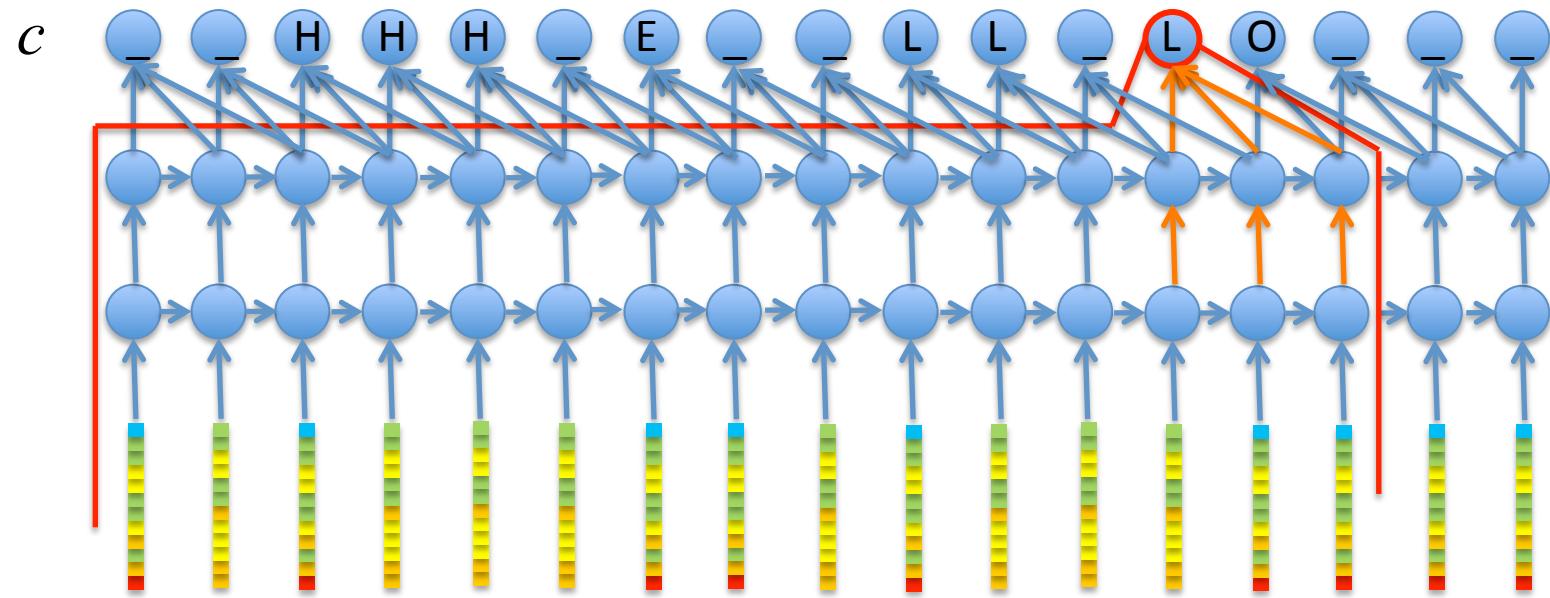
- Fix: bake limited context into model structure.



- Caveat: May need to compute upper layers quickly after sufficient context arrives. Can be easier if context is near top.

# Latency

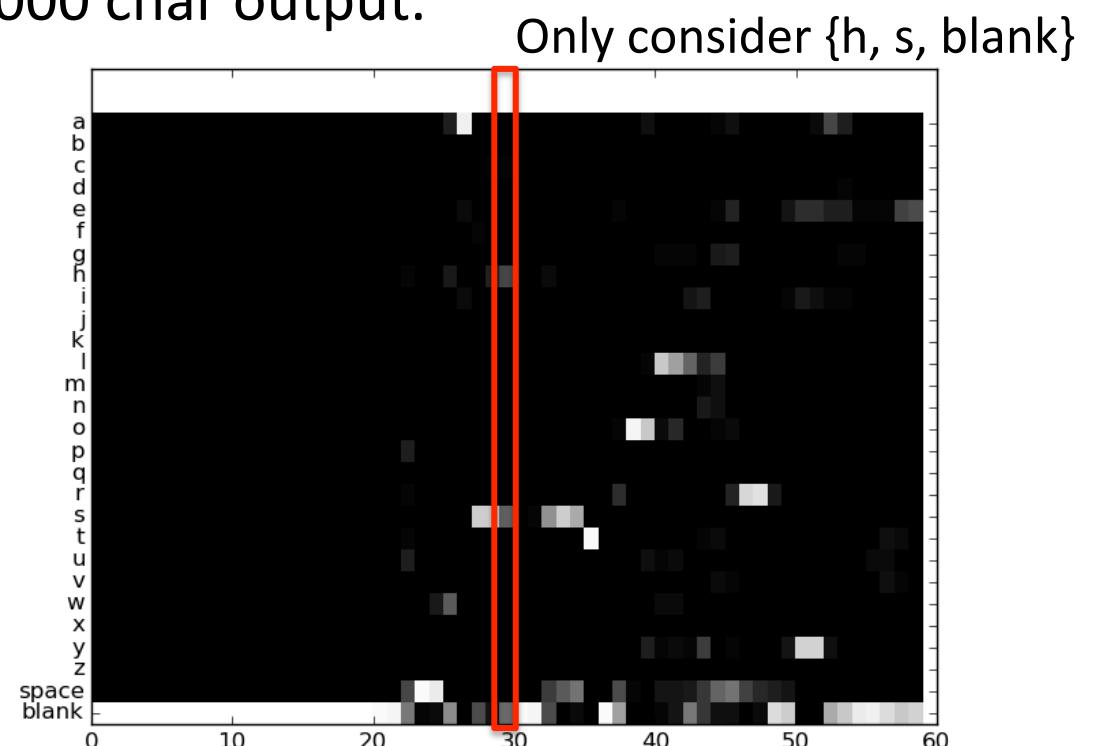
- Fix: Move most of context to the top.



- Can easily compute/recompute top layers online.

# Pruning candidates

- For models with many character outputs (e.g., Mandarin), decoding slows down.
  - Trick: only consider top 99% of characters according to CTC.
  - 150x speedup for 5000 char output.



# Throughput

- Large DNN/RNN models hard to deploy on CPUs.
- Large DNN/RNN models run great on GPUs.
  - But only if “batch size” is high enough.
  - Processing 1 audio stream at a time is inefficient.

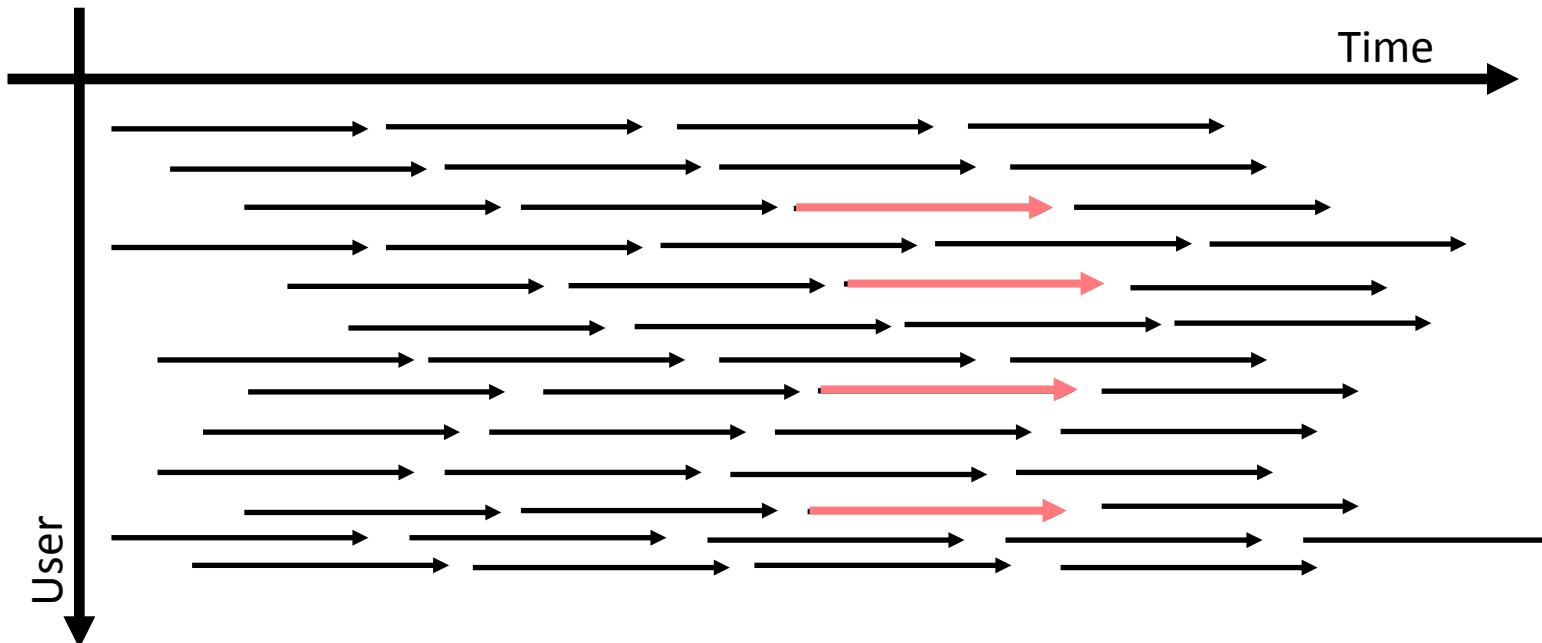
Performance for K1200 GPU:

Batch Size	FLOPS	Throughput
1	0.065 TFLOPs	1x
10	0.31 TFLOPs	5x
32	0.92 TFLOPs	14x

[Chris Fougner]

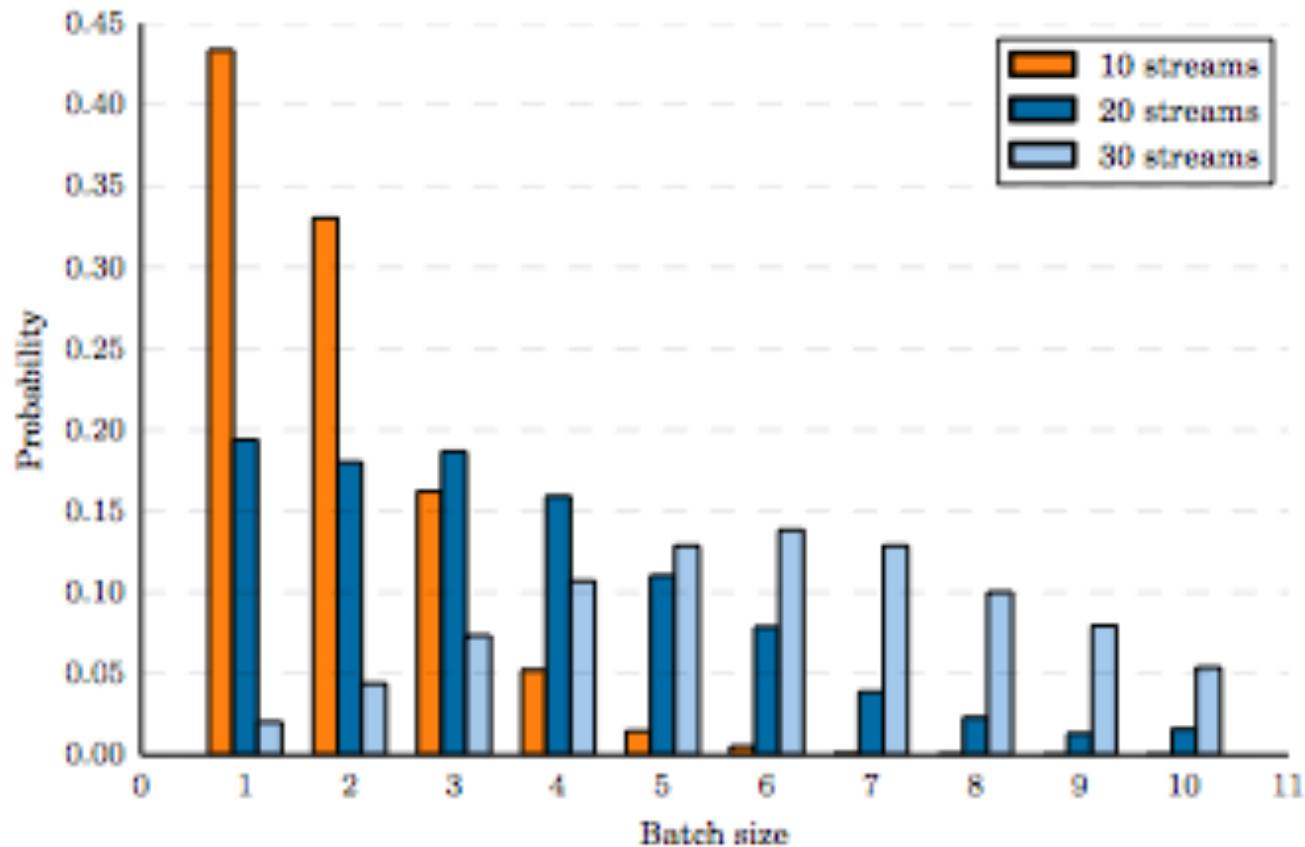
# Throughput

- Batch packets together as data comes in.



- Arrows represent packet of speech data (e.g., 100ms).
- Idea: Process packets that arrive at similar times in parallel.

# Throughput



With ~30 concurrent users, few GEMM batches have less than size 4.

# Summary

- Deep Learning makes first steps to state-of-art speech system simpler than ever.
- Performance significantly driven by data & models.
  - Focus on scaling data + compute.
  - Try more models, make more progress!
- Mature enough for production.
  - DeepSpeech model is live in Mandarin & English.

**Starter code:** [github.com/baidu-research/ba-dls-deepspeech](https://github.com/baidu-research/ba-dls-deepspeech)

# References

- Gales and Young. "The Application of Hidden Markov Models in Speech Recognition" Foundations and Trends in Signal Processing, 2008.
- Jurafsky and Martin. "Speech and Language Processing". Prentice Hall, 2000.
- Dzmitry Bahdanau, Jan Chorowski, Dmitriy Serdyuk, Philemon Brakel, Yoshua Bengio. "End-to-End Attention-based Large Vocabulary Speech Recognition." 2015. arxiv.org/abs/1508.04395
- Bourlard and Morgan. "CONNECTIONIST SPEECH RECOGNITION: A Hybrid Approach". Kluwer Publishing, 1994.
- William Chan, Navdeep Jaitly, Quoc V. Le, Oriol Vinyals. "Listen Attend Spell" 2015. arxiv.org/abs/1508.01211
- Jianmin Chen, Rajat Monga, Samy Bengio, Rafal Jozefowicz, "Revisiting Distributed Synchronous SGD." ICLR 2016 arXiv:1604.00981
- A Graves, S Fernández, F Gomez, J Schmidhuber. "Connectionist temporal classification: labelling unsegmented sequence data with recurrent neural networks." ICML, 2006.
- Dahl, Yu, Deng, Acero, "Large Vocabulary Continuous Speech Recognition with Context-Dependent DBN-HMMs". ICASSP, 2011
- Dean, J., Corrado, G., Monga, R., Chen, K., Devin, M., Mao, M., Senior, A., Tucker, P., Yang, K., Le, Q.V. and Ng, A.Y. "Large scale distributed deep networks." NIPS 2012
- Hannun, Maas, Jurafsky, Ng. "First-Pass Large Vocabulary Continuous Speech Recognition using Bi-Directional Recurrent DNNs" ArXiv: 1408.2873
- Hannun, et al. "Deep Speech: Scaling up end-to-end speech recognition". ArXiv:1412.5567
- H. Hermansky, "Perceptual linear predictive (PLP) analysis of speech", J. Acoust. Soc. Am., vol. 87, no. 4, pp. 1738-1752, Apr. 1990.
- H. Hermansky and N. Morgan, "RASTA processing of speech", IEEE Trans. on Speech and Audio Proc., vol. 2, no. 4, pp. 578-589, Oct. 1994.
- Vassil Panayotov, Guoguo Chen, Daniel Povey and Sanjeev Khudanpur, "LibriSpeech: an ASR corpus based on public domain audio books." ICASSP 2015
- H. Schwenk, "Continuous space language models", 2007.